

LED S ML

AS O LE 2

CE LM GO

SE MS 1 2

EG O C O 1

Recurrent Neural Networks and Long Short-Term Memory

Lluís Guasch & Carlos Cueto

GTAs online:
Alex
George

1. Recurrent Neural Networks (RNNs)
2. Long Short-Term Memory (LSTM) networks
3. Text representation

1. Recurrent Neural Networks (RNNs)
2. Long Short-Term Memory (LSTM) networks
3. Text representation


Examples of sequential data:

- ▶ Text considered as a sequence of words or characters
- ▶ Continuous parameter which is a function of time (e.g. stock price)
- ▶ Sequence of images in a video-clip
- ▶ Sequence of labels in a genome sequence
- ▶ Vertical sequences of lithologies in the subsurface
- ▶ ... plus many others

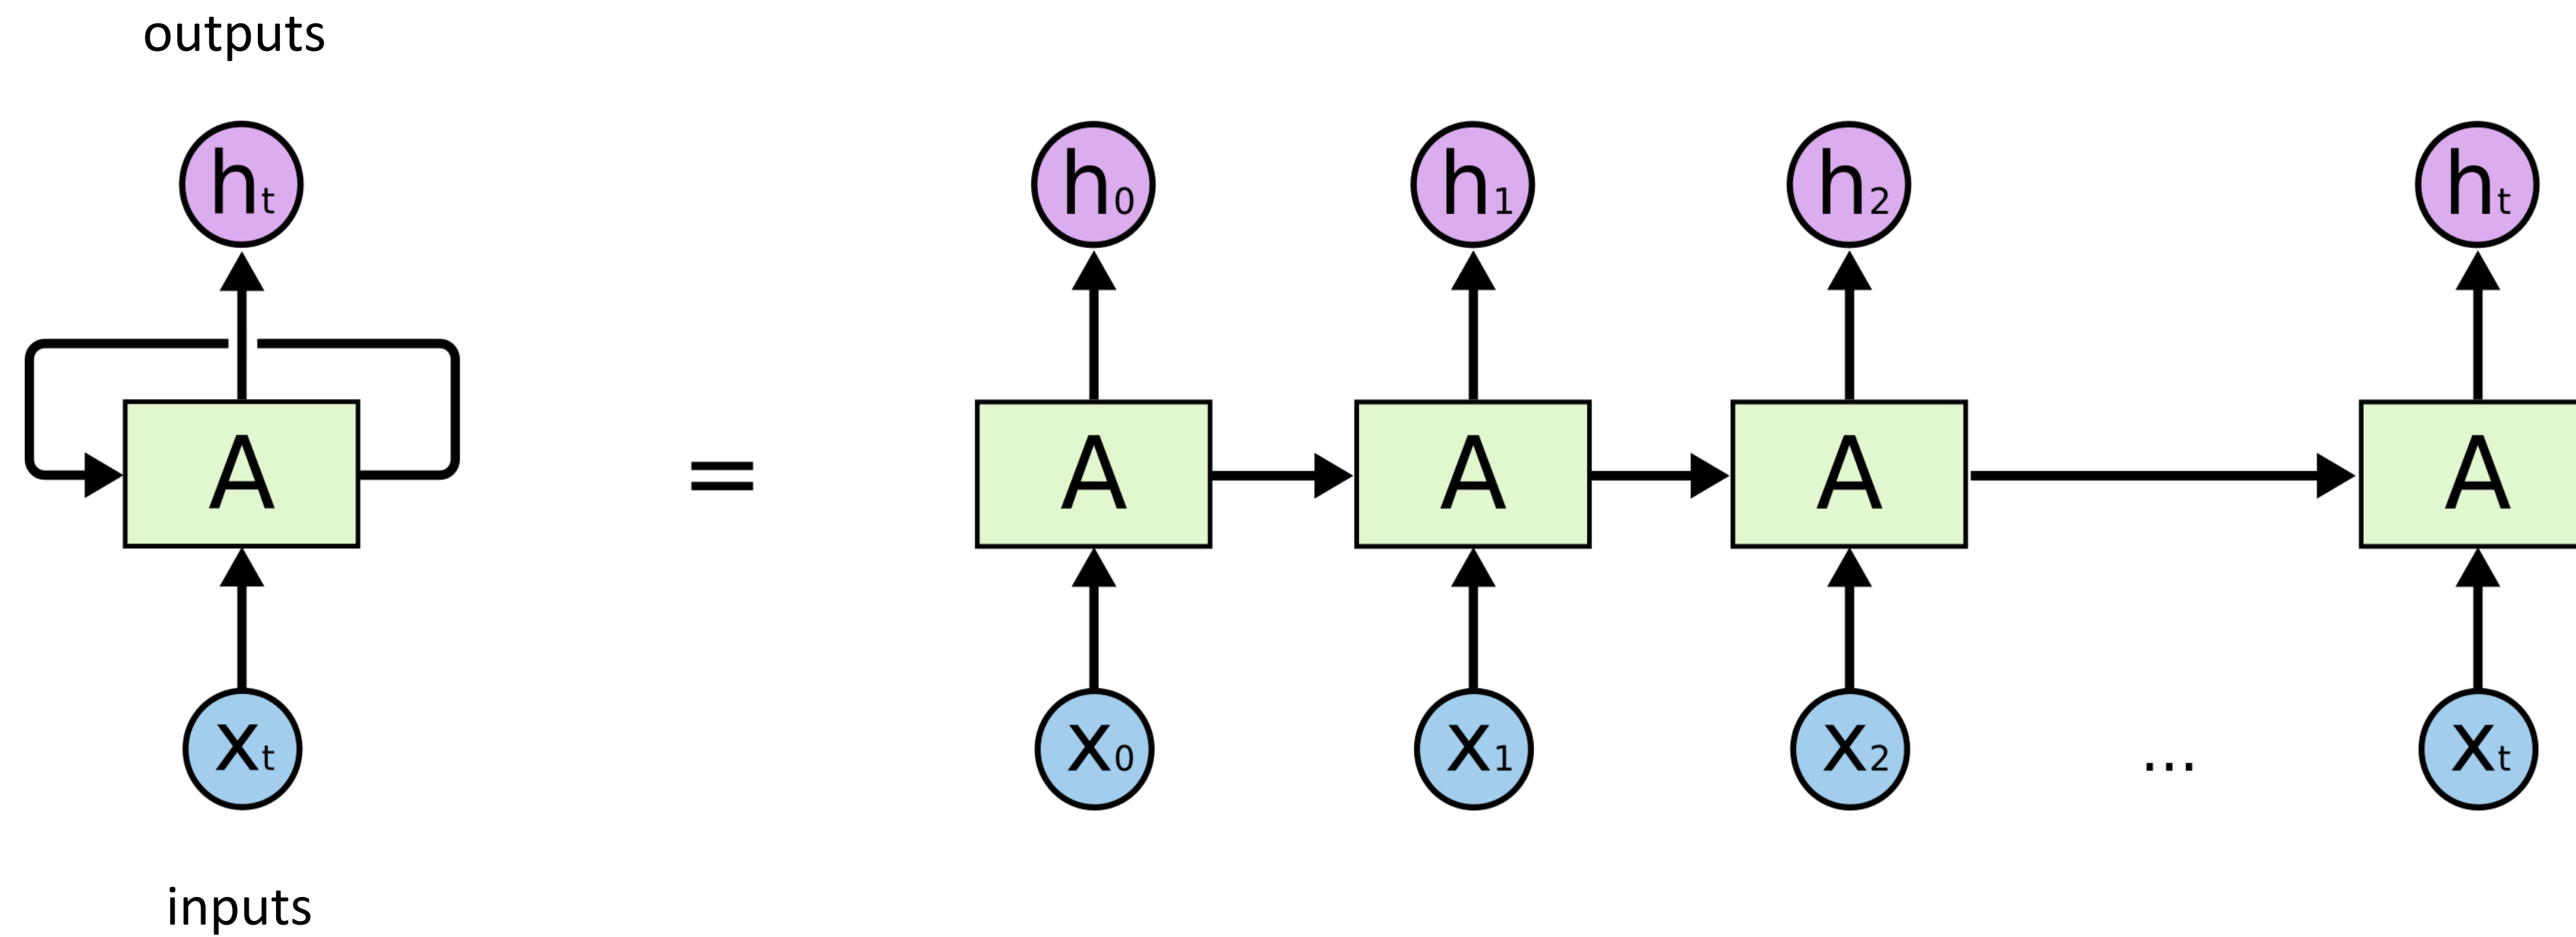
Applications of RNNs to sequential data:

- ▶ Text generation “in the style of”
- ▶ Provide “sentiment analysis” on a piece of text
- ▶ Automatic “speech-to-text” as in MS Teams
- ▶ Automatic foreign language translation
- ▶ Prediction of stock price on the basis of historical data
- ▶ Label sequence of images in a video-clip

Basic structure of RNNs:



Basic structure of RNNs:



Simple (unrealistic) example:

Example:

Based on a (usually long) text used for training, we wish to generate “similar” text on a character by character basis.

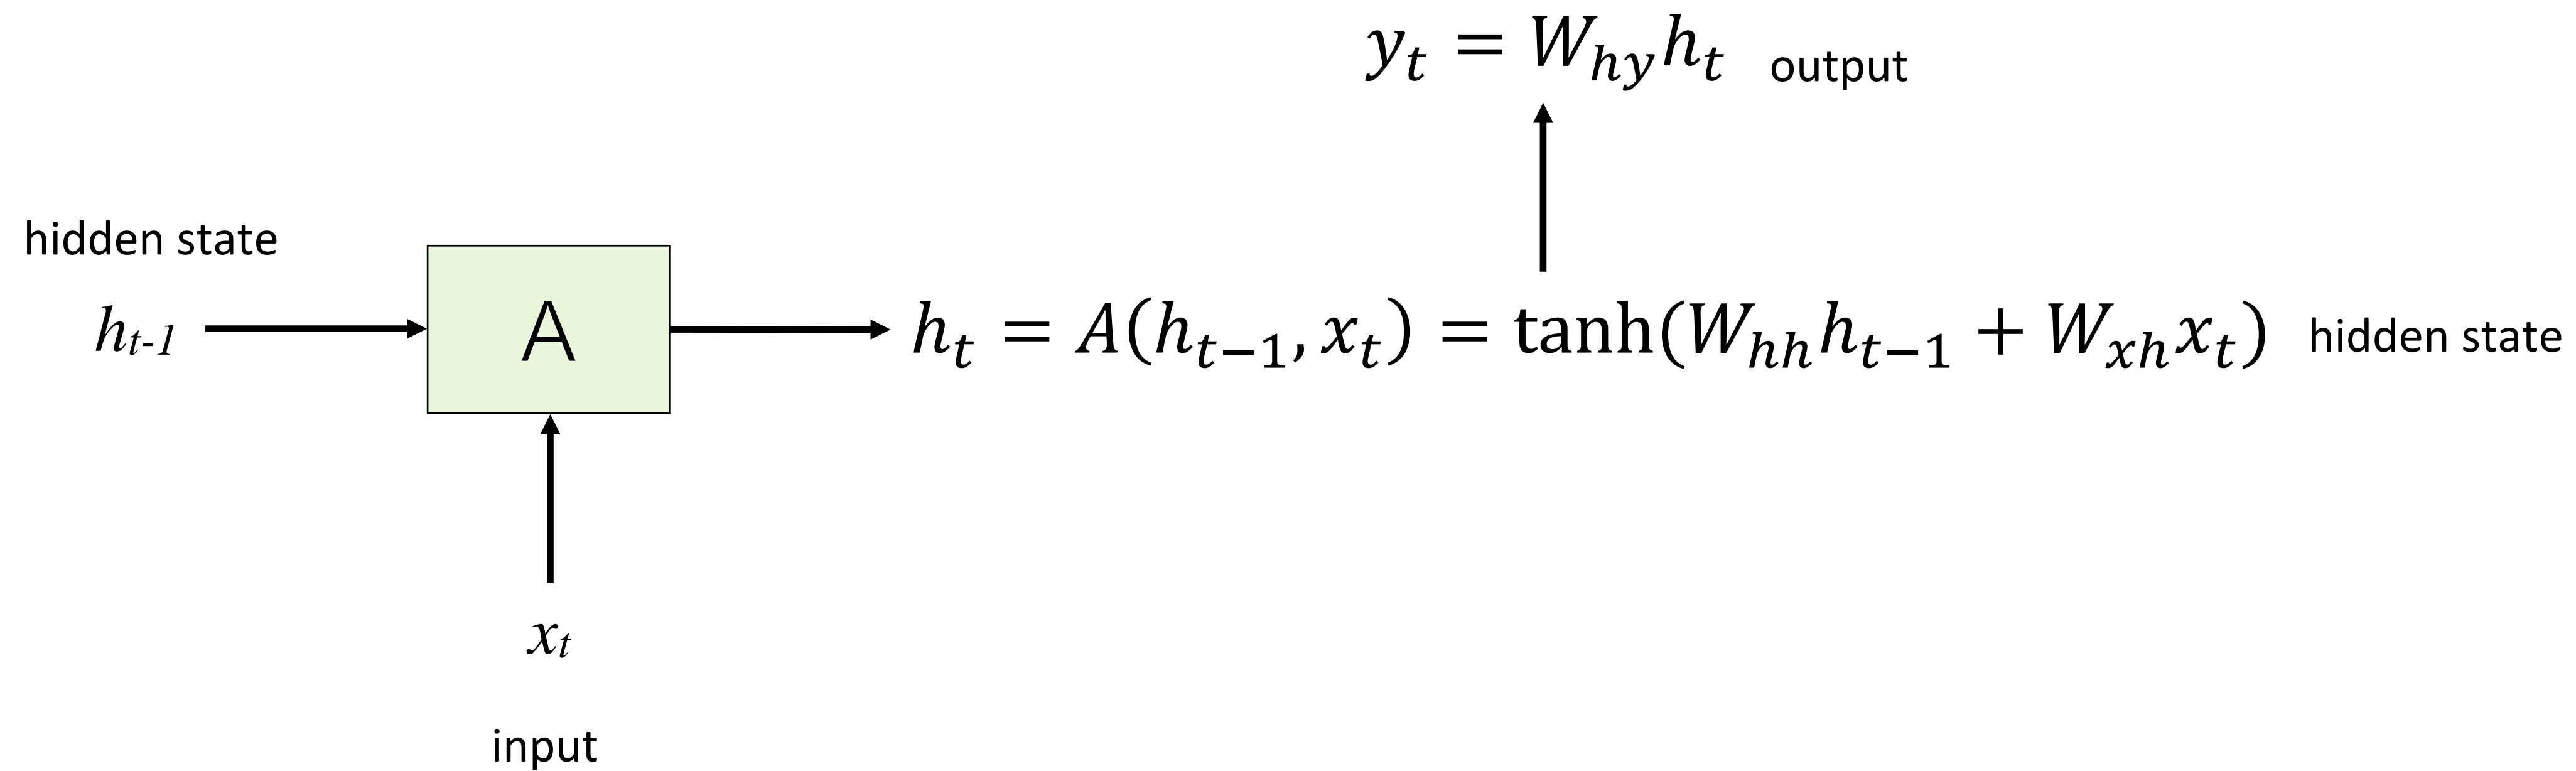
Given a training string: `Hello world!`

We identify the dictionary associated with the input (ignoring capitalisation) contains 9 characters: (d, e, h, l, o, r, w, , !)

And we hot-encode the dictionary:

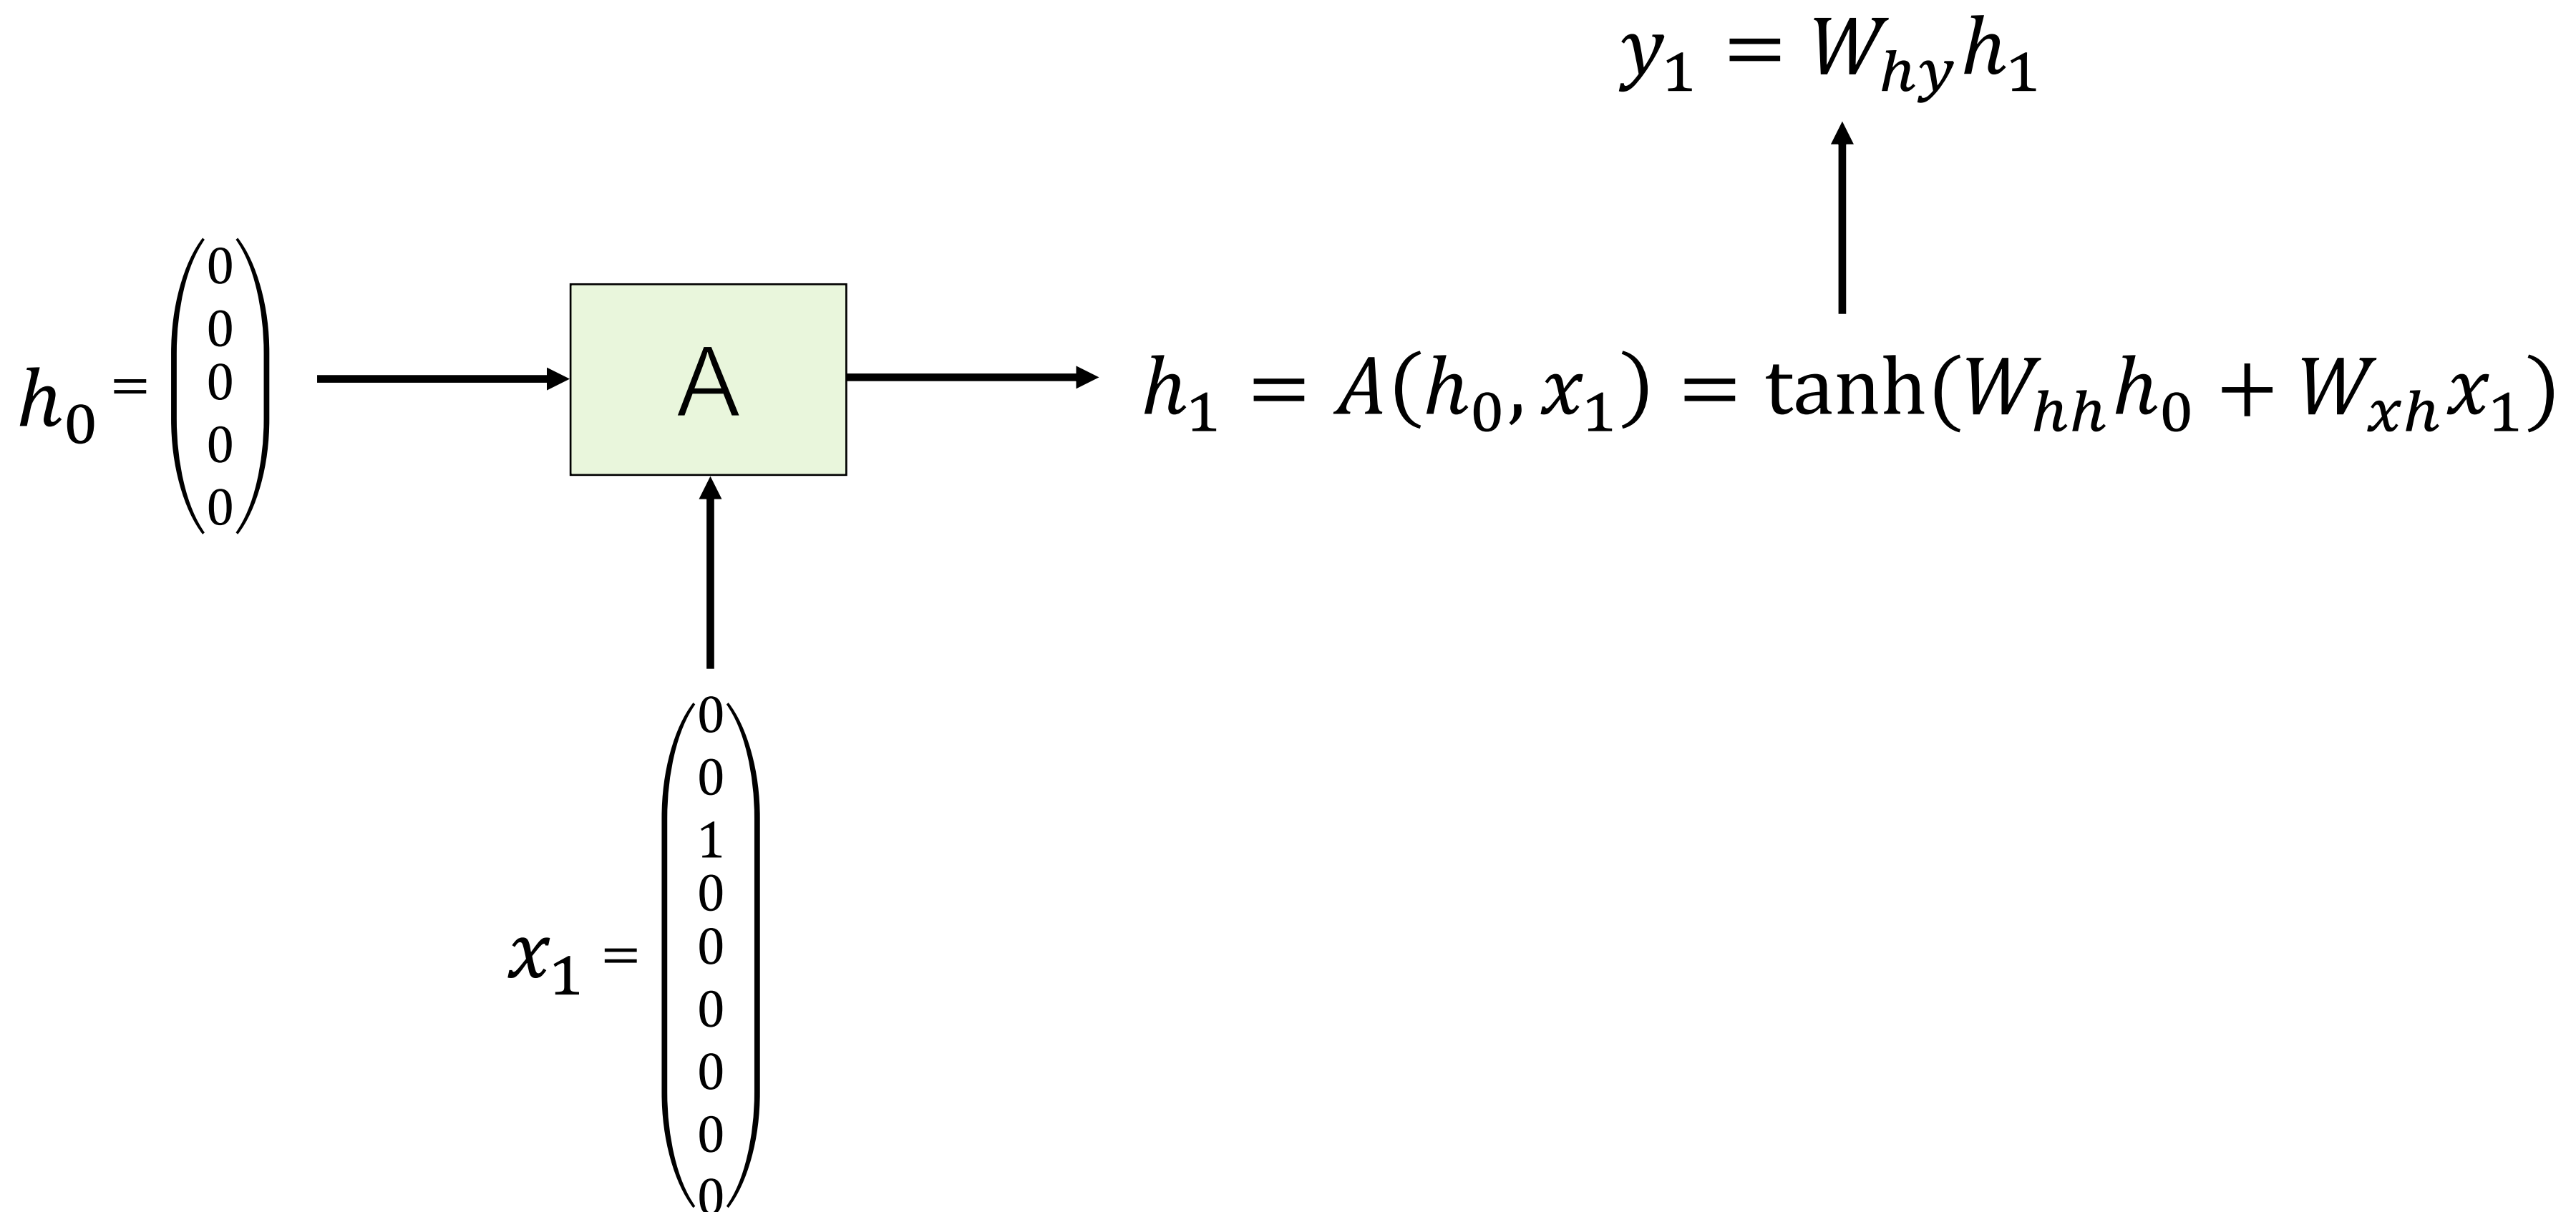
$$h = \begin{pmatrix} 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix} \quad ! = \begin{pmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \end{pmatrix} \quad \dots$$

How does a basic RNN cell work?




How does training work? Back to our simple example:

We define the size of our hidden state tensor h (hyperparameter) to be 5, and start with our first input character ('h'):



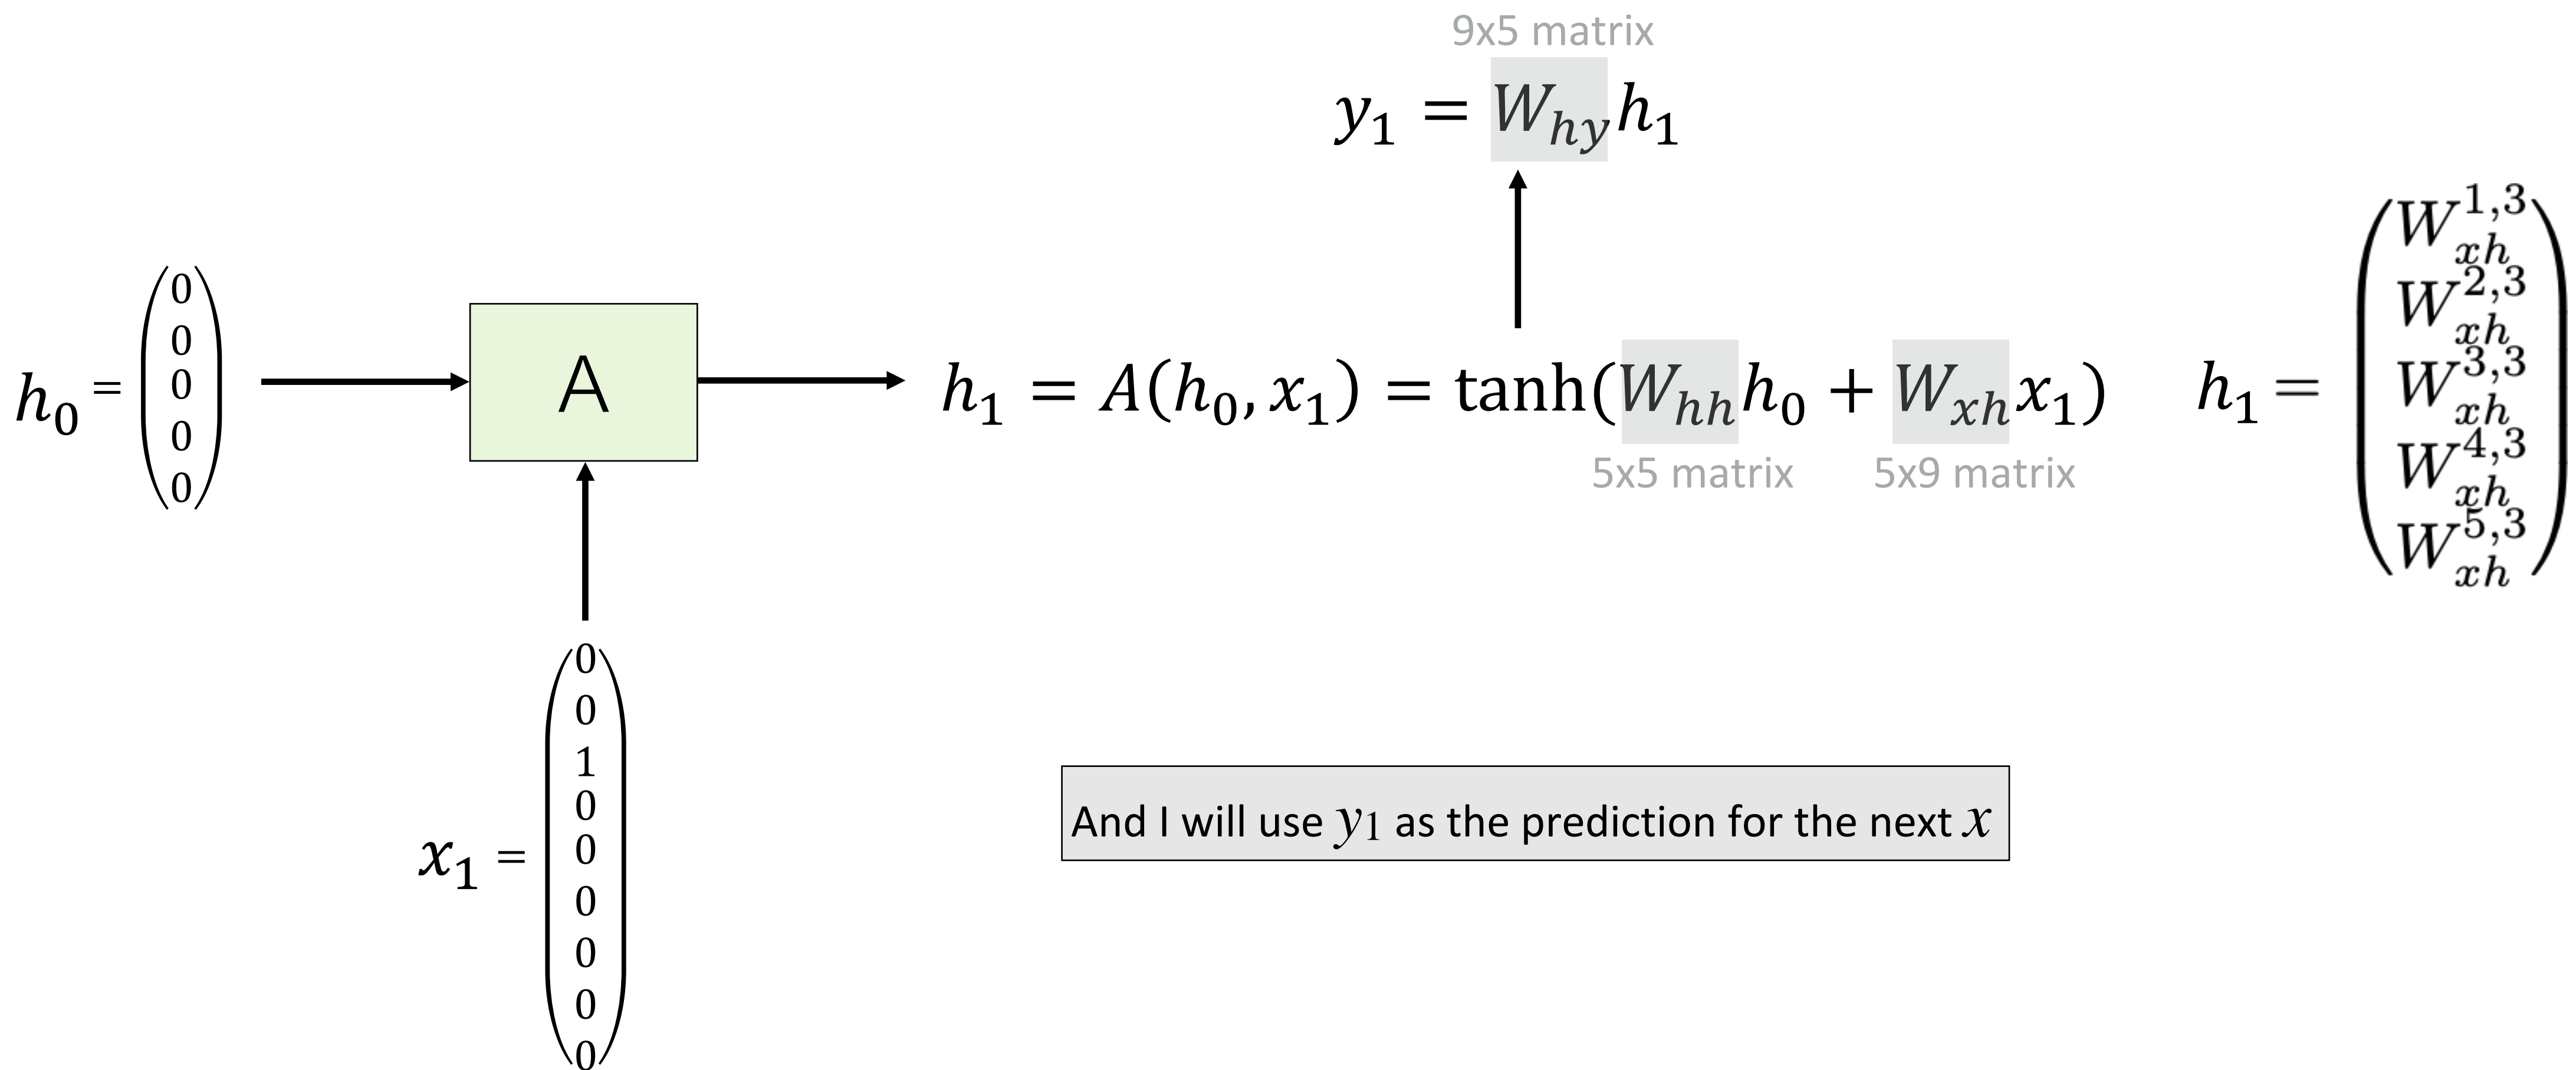
How does training work? Back to our simple example:

We define the size of our hidden state tensor h (hyperparameter) to be 5, and start with our first input character ('h'):



How does training work? Back to our simple example:

We define the size of our hidden state tensor h (hyperparameter) to be 5, and start with our first input character ('h'):



How does training work? The loss function:

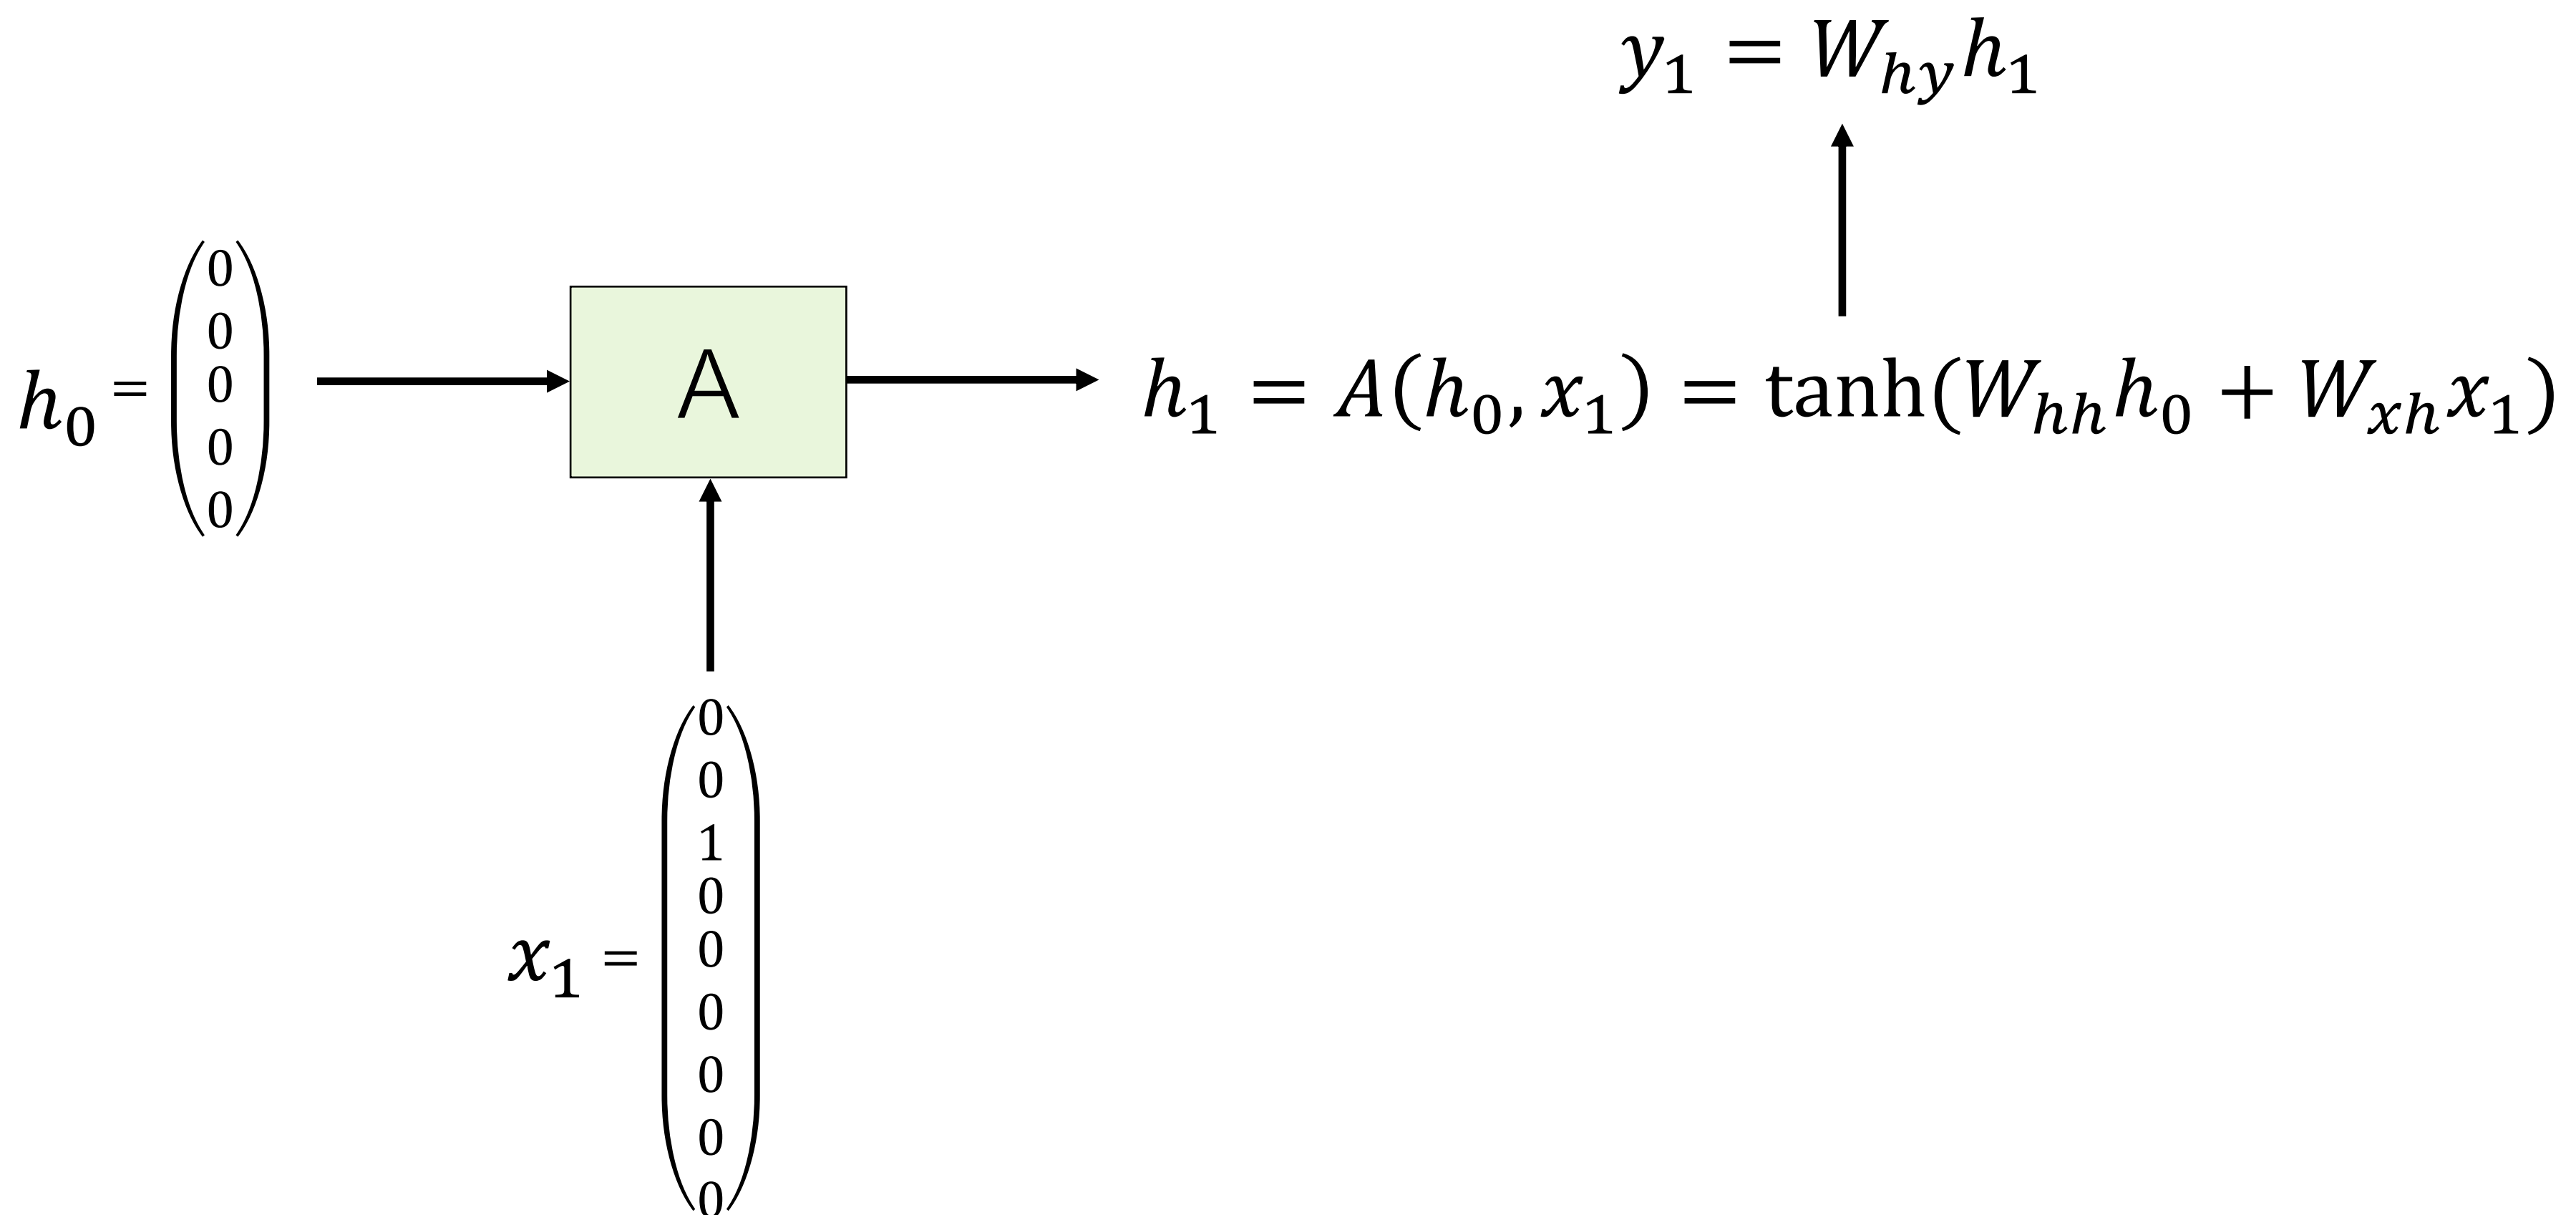
$loss =$

$$y_1 = W_{hy}h_1 \longrightarrow \textit{Cross-entropy} (\textit{Softmax} (y_1) , x_2)$$

Loss function corresponding to
the first predicted letter

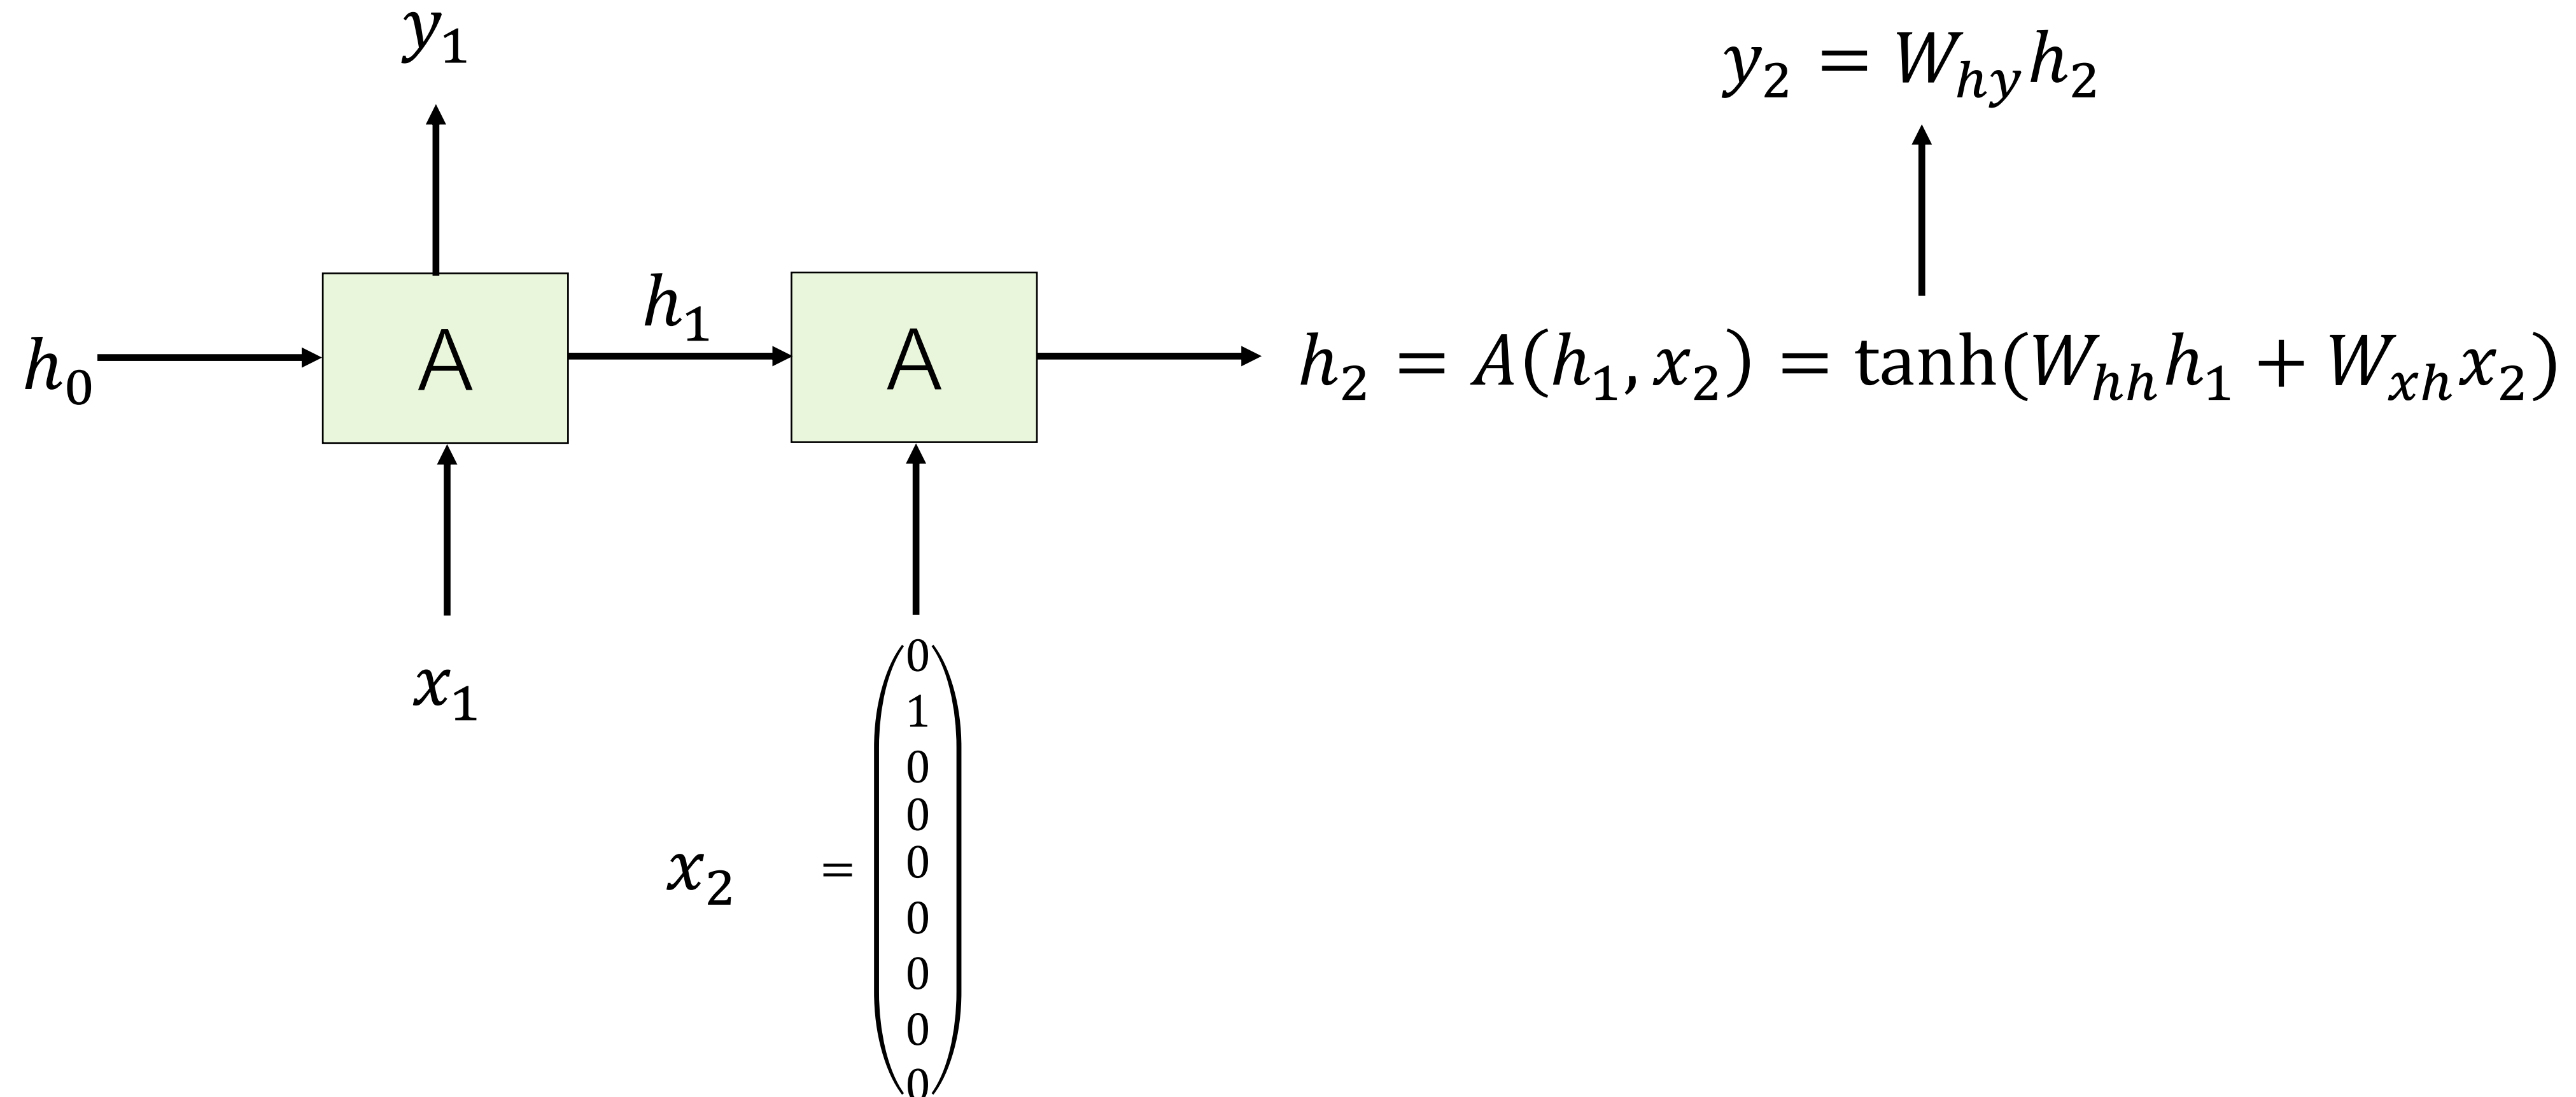
remember: hot-
encoded

How does training work for the next characters? Back to our previous example:



How does training work for the next characters? Back to our previous example:

We concatenate the **same** cell after the output of our previous one:



How does training work? The loss function:

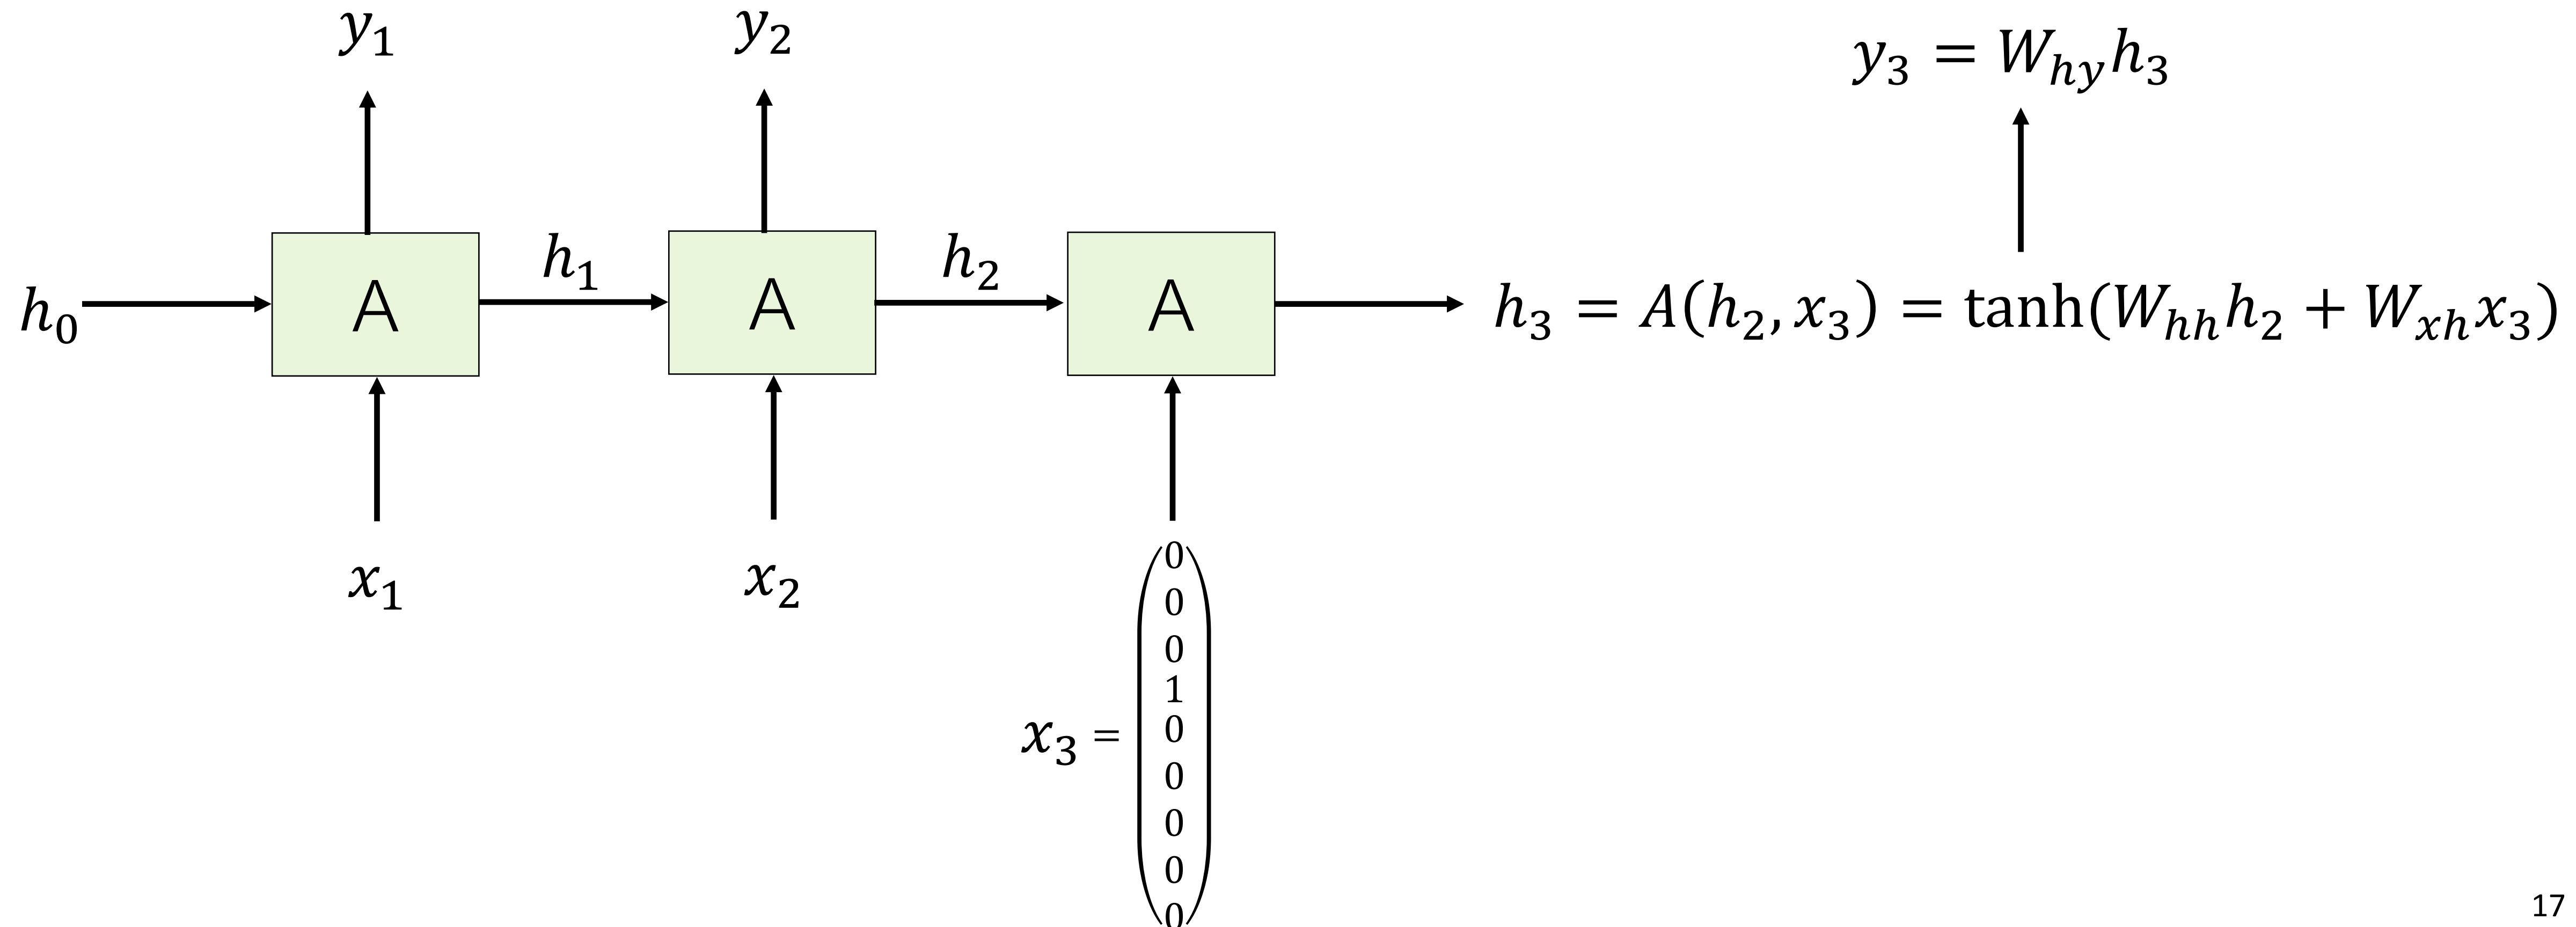
$$\begin{aligned} & \text{loss} = \\ & y_1 = W_{hy}h_1 \longrightarrow \text{Cross-entropy} (\text{Softmax} (y_1) , x_2) \\ & \quad + \\ & \quad \text{Cross-entropy} (\text{Softmax} (y_2) , x_3) \end{aligned}$$

Loss function corresponding to
the second predicted letter

remember: hot-
encoded

How does training work for the next characters? Back to our previous example:

We concatenate the **same** cell after the output of our previous one:



How does training work? The loss function:

$$\begin{aligned}
 & \text{loss} = \\
 & y_1 = W_{hy}h_1 \longrightarrow \text{Cross-entropy} (\text{Softmax} (y_1) , x_2) \\
 & \quad + \\
 & \quad \text{Cross-entropy} (\text{Softmax} (y_2) , x_3) \\
 & \quad + \\
 & \quad \text{Cross-entropy} (\text{Softmax} (y_3) , x_4) \\
 & \quad + \dots
 \end{aligned}$$

We repeat the process for all the characters in our input and add their loss contributions

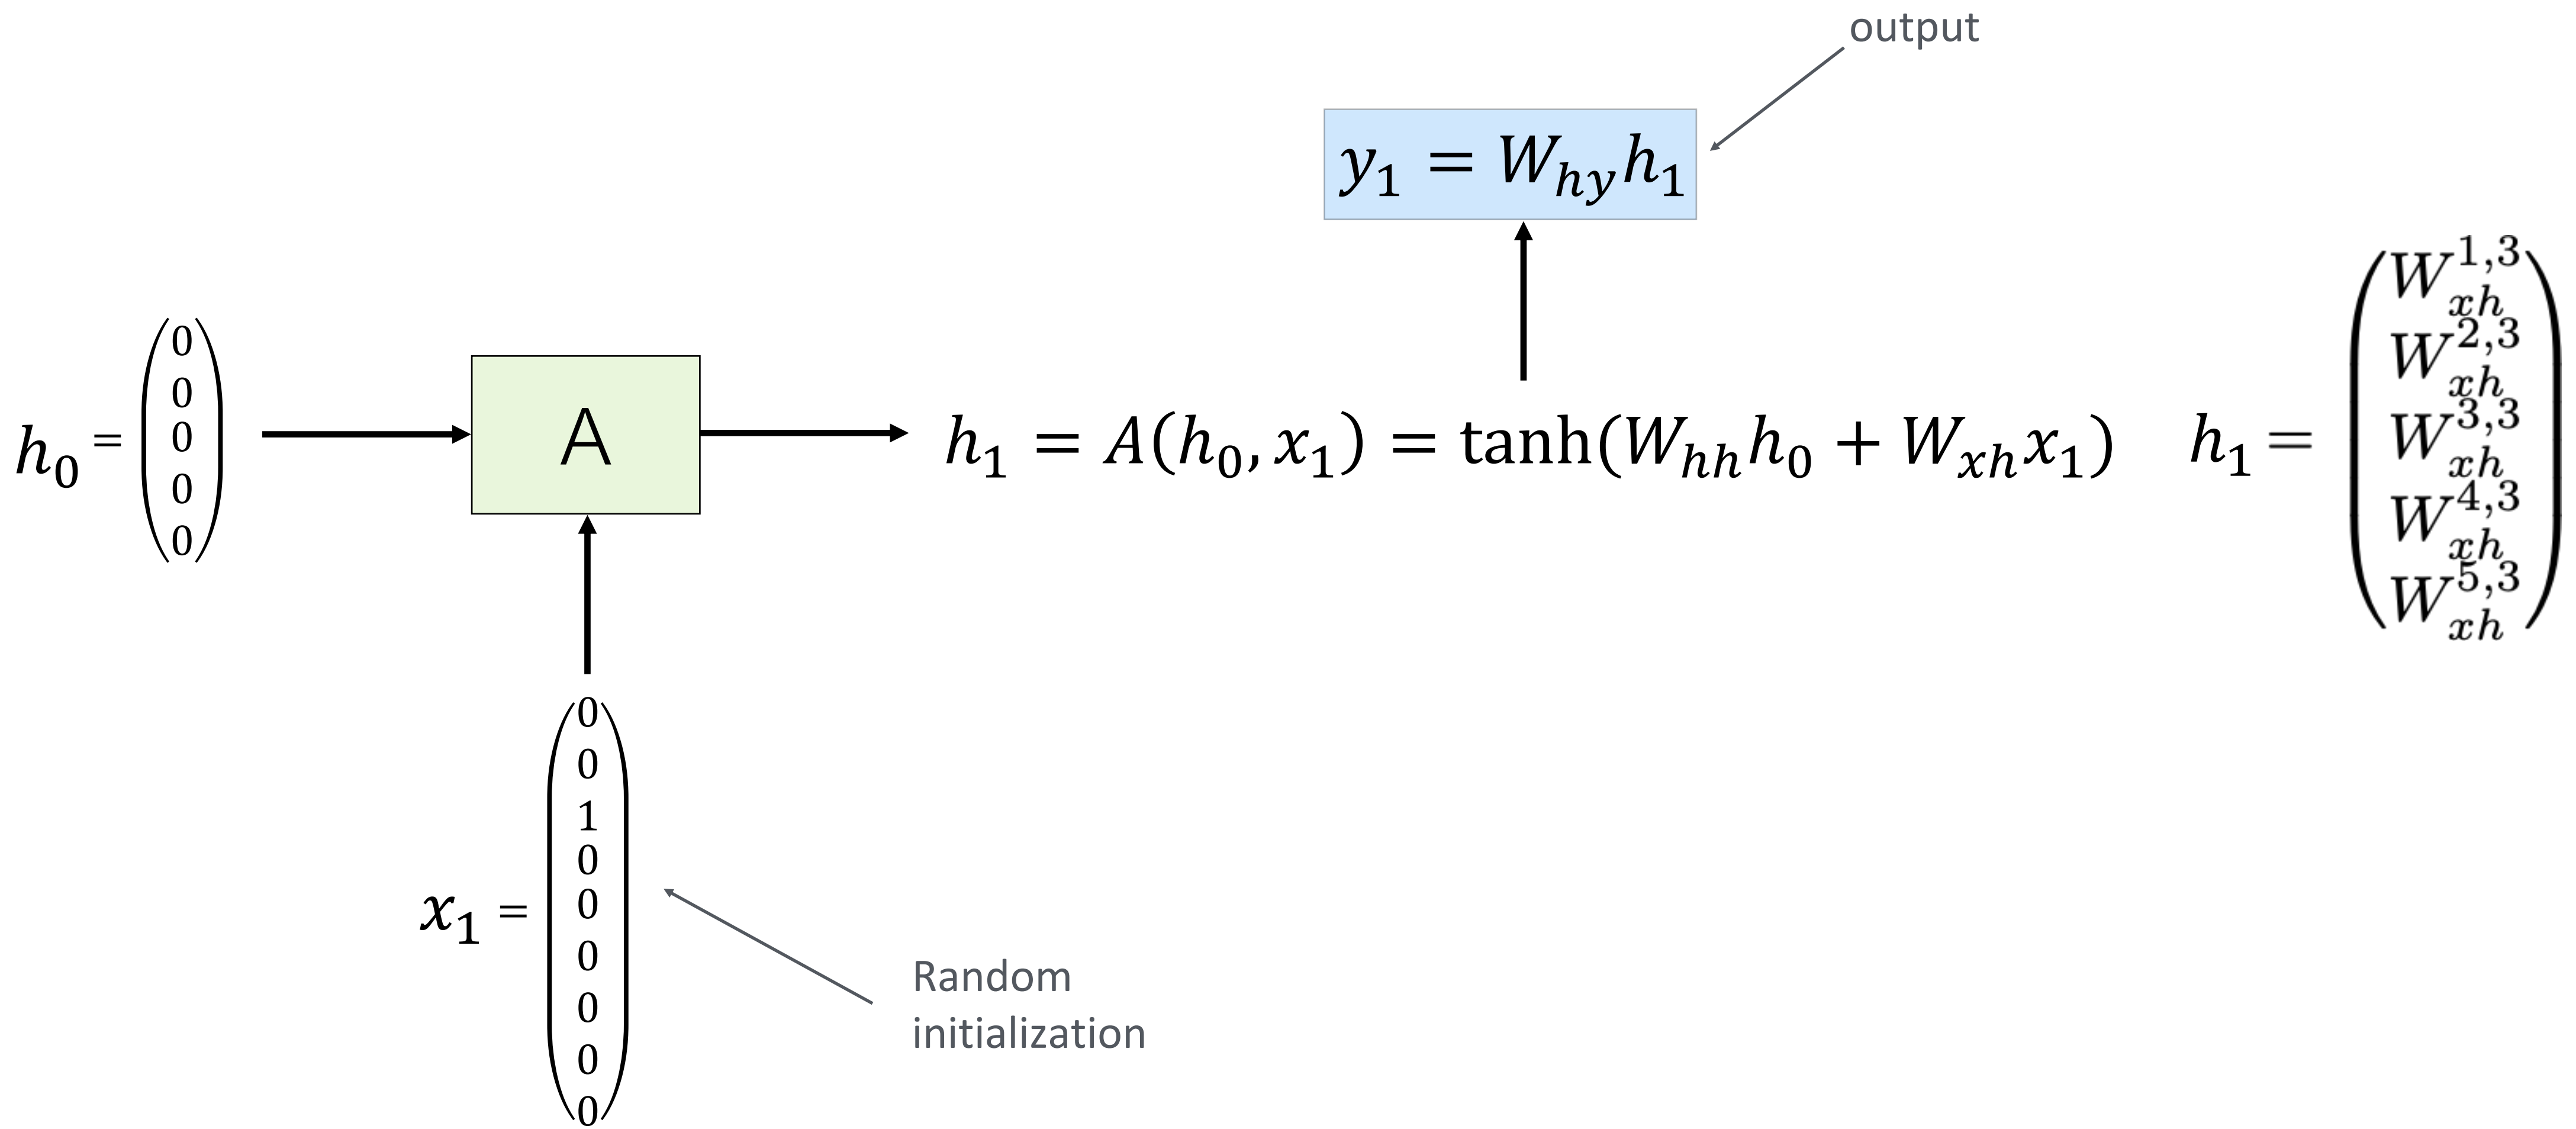
Loss function corresponding to the third predicted letter

remember: hot-encoded

And then use gradient descent to train our weights in the matrices


Text generation using a trained network

Now we are going to use our trained RNN to generate some text:




Text generation using a trained network

Now we are going to use our trained RNN to generate some text:



Text generation using a trained network

Now we are going to use our trained RNN to generate some text:



Recap of training an RNN

1. Take first character x_1 , combine with current hidden state h_0 (its dimension is a hyperparameter) to derive h_1 , calculate associated output vector y_1 (of size equal to vocabulary), transform y_1 into Softmax vector, calculate cross-entropy with second character x_2 of the sequence.
2. Take the second character as input to repeat the step 1. and calculate loss function again using the value of the third character of the sequence, and so on until we reach the end of the sequence.
3. Add up all the above loss functions.
4. Do backpropagation to calculate gradients according to each trainable parameter.
5. Modify parameters by gradient descent.
6. Move to next sequence of p characters, until end of training set is reached. Size of p is a hyperparameter too!


Recap of generating text with an RNN

Generate a new sequence of n characters:

1. Randomly sample first character x_1 .
2. Combine with initial (zero) hidden state vector value h_0 to derive new vector h_1 .
3. Calculate associated output vector y_1 of size equal to vocabulary size.
4. Transform y_1 into Softmax vector.
5. Sample new character x_2 based on Softmax probability.
6. Continue until the required number of characters have been generated.

If we were using words instead of characters, the dictionary would be much bigger

Synthetic notation for RNNs:



Simple implementation in python:

<https://gist.github.com/karpathy/d4dee566867f8291f086>

And after a few hours of training, we start to get some decent writing:

```
VIOLA:  
Why, Salisbury must find his flesh and thought  
That which I am not aps, not a man and in fire,  
To show the reining of the raven and the wars  
To grace my hand reproach within, and not a fair are hand,  
That Caesar and my goodly father's world;  
When I was heaven of presence and our fleets,  
We spare with hours, but cut thy council I am great,  
Murdered and by thy master's ready there  
My power to give thee but so much as hell:  
Some service in the noble bondman here,  
Would show him to her wine.
```

```
KING LEAR:  
O, if you were a feeble sight, the courtesy of your law,  
Your sight and several breath, will wear the gods  
With his heads, and my hands are wonder'd at the deeds,  
So drop upon your lordship's head, and your opinion  
Shall be against your honour.
```


Analysis of the role of h

Cell sensitive to position in line:

The sole importance of the crossing of the Berezina lies in the fact that it plainly and indubitably proved the fallacy of all the plans for cutting off the enemy's retreat and the soundness of the only possible line of action--the one Kutuzov and the general mass of the army demanded--namely, simply to follow the enemy up. The French crowd fled at a continually increasing speed and all its energy was directed to reaching its goal. It fled like a wounded animal and it was impossible to block its path. This was shown not so much by the arrangements it made for crossing as by what took place at the bridges. When the bridges broke down, unarmed soldiers, people from Moscow and women with children who were with the French transport, all--carried on by vis inertiae--pressed forward into boats and into the ice-covered water and did not, surrender.

Cell that turns on inside quotes:

"You mean to imply that I have nothing to eat out of.... On the contrary, I can supply you with everything even if you want to give dinner parties," warmly replied Chichagov, who tried by every word he spoke to prove his own rectitude and therefore imagined Kutuzov to be animated by the same desire.

Kutuzov, shrugging his shoulders, replied with his subtle penetrating smile: "I meant merely to say what I said."

Cell that robustly activates inside if statements:

```
static int __dequeue_signal(struct sigpending *pending, sigset_t *mask,
                           siginfo_t *info)
{
    int sig = next_signal(pending, mask);
    if (sig) {
        if (current->notifier) {
            if (sigismember(current->notifier_mask, sig)) {
                if (!(current->notifier)(current->notifier_data)) {
                    clear_thread_flag(TIF_SIGPENDING);
                    return 0;
                }
            }
        }
        collect_signal(sig, pending, info);
    }
    return sig;
}
```

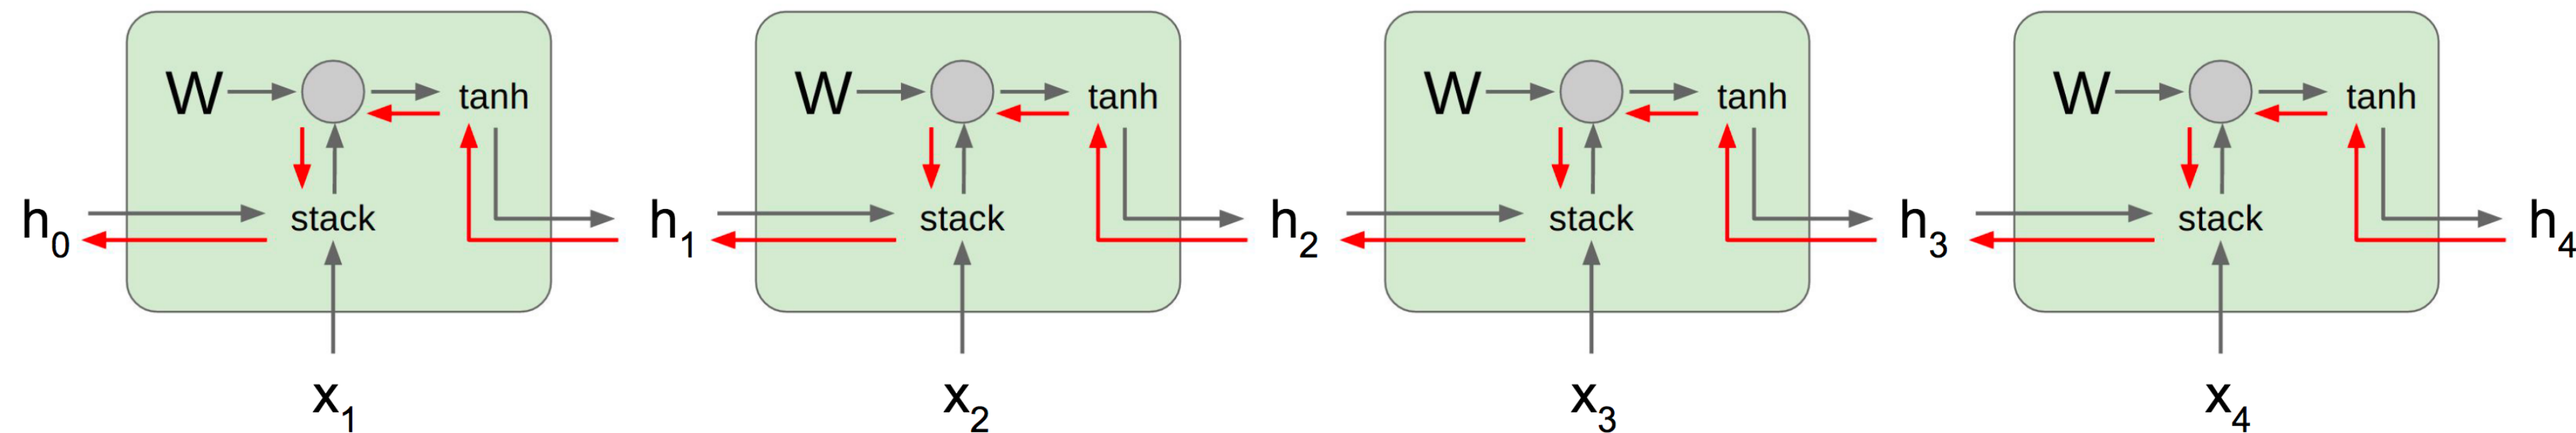

Analysis of the role of h

A large portion of cells are not easily interpretable. Here is a typical example:

```
/* Unpack a filter field's string representation from user-space
 * buffer. */
char *audit_unpack_string(void **bufp, size_t *remain, size_t len)
{
    char *str;
    if (!*bufp || (len == 0) || (len > *remain))
        return ERR_PTR(-EINVAL);
    /* Of the currently implemented string fields, PATH_MAX
     * defines the longest valid length.
     */
}
```

In their study, Karpathy et al found that only ~5% of the h positions were interpretable

Backpropagation through RNNs



Computing gradients involves many factors of W and repeated \tanh activations and that creates problems:

- ▶ Exploding gradients: can be solved by gradient clipping
- ▶ Vanishing gradients: requires changes in the RNN architecture

the derivative of \tanh is in the range $(0, 1)$ and we apply it multiple times using the chain rule in backprop

Issues with RNNs

- Exploding and vanishing gradients
- Lack of long term memory

Issues with RNNs



- Exploding and vanishing gradients
- Lack of long term memory





LSTMs to the rescue!

1. Recurrent Neural Networks (RNNs)
2. Long Short-Term Memory (LSTM) networks
3. Text representation

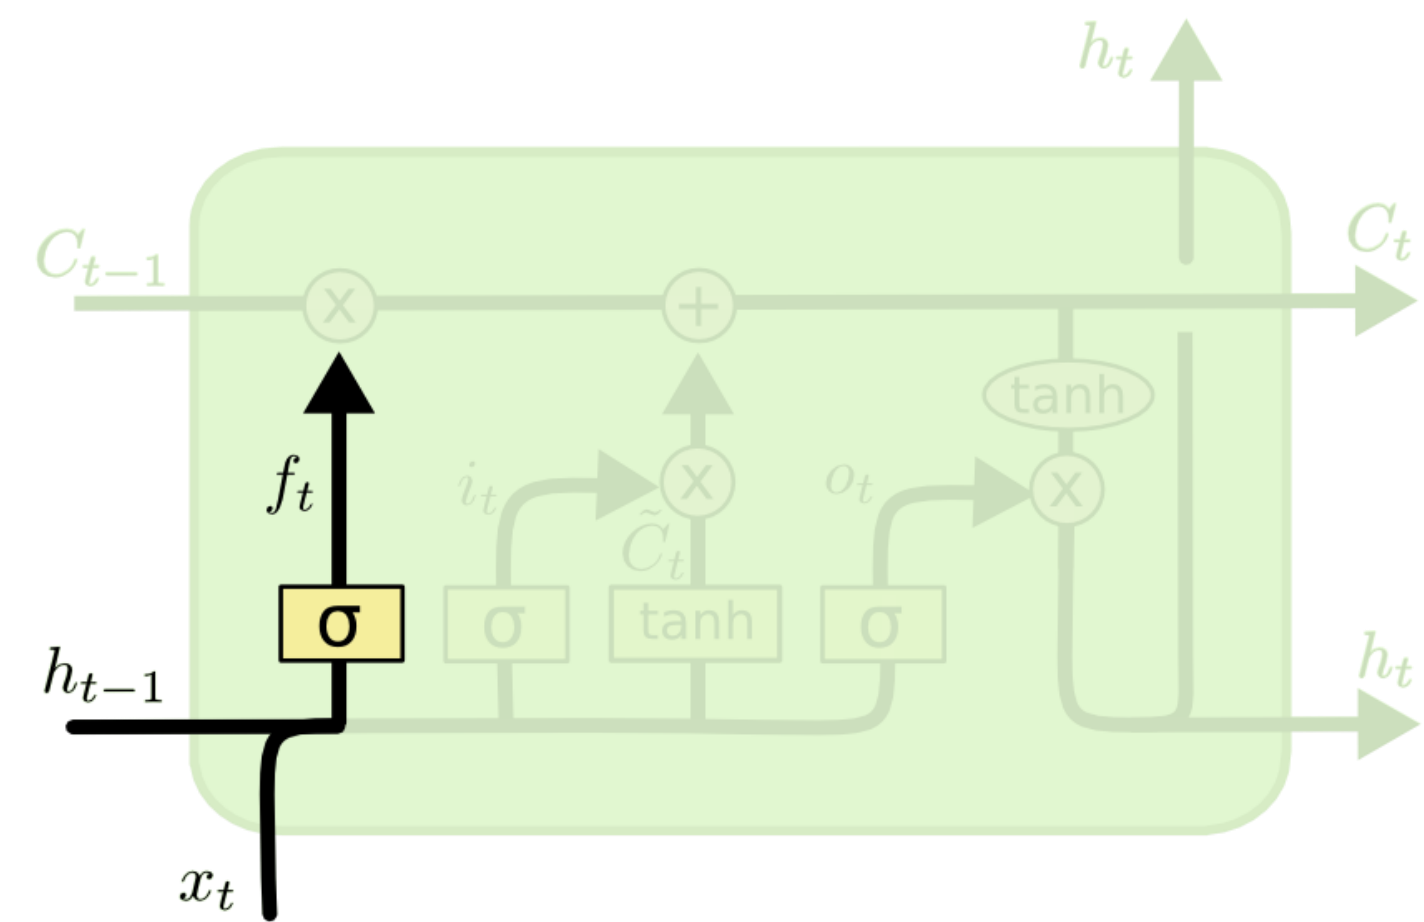
Basic structure of an LSTM:



The core idea behind LSTMs: the **cell state** directly connects the whole chain of cells.




The **forget gate**: is there any information in the cell state that we no longer need?

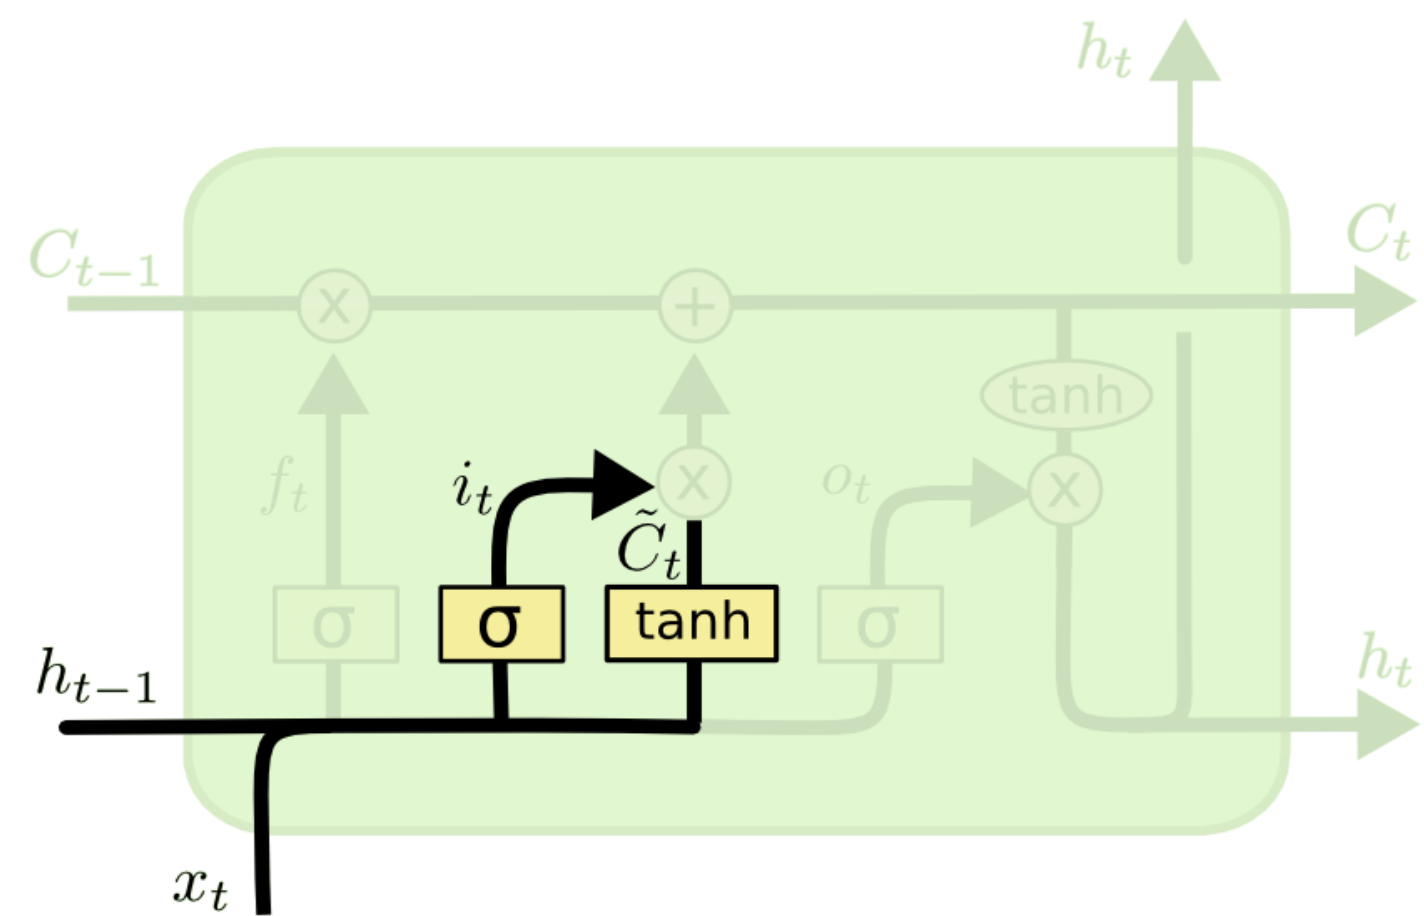


sigmoid normalises between 0 and 1

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$




The **input gate** and the **candidate update**: what part of the input should we use to update the cell state?




$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$
$$\tilde{C}_t = \tanh(W_c \cdot [h_{t-1}, x_t] + b_c)$$

tanh normalises between -1 and 1,
like cell state




Updating the cell state:



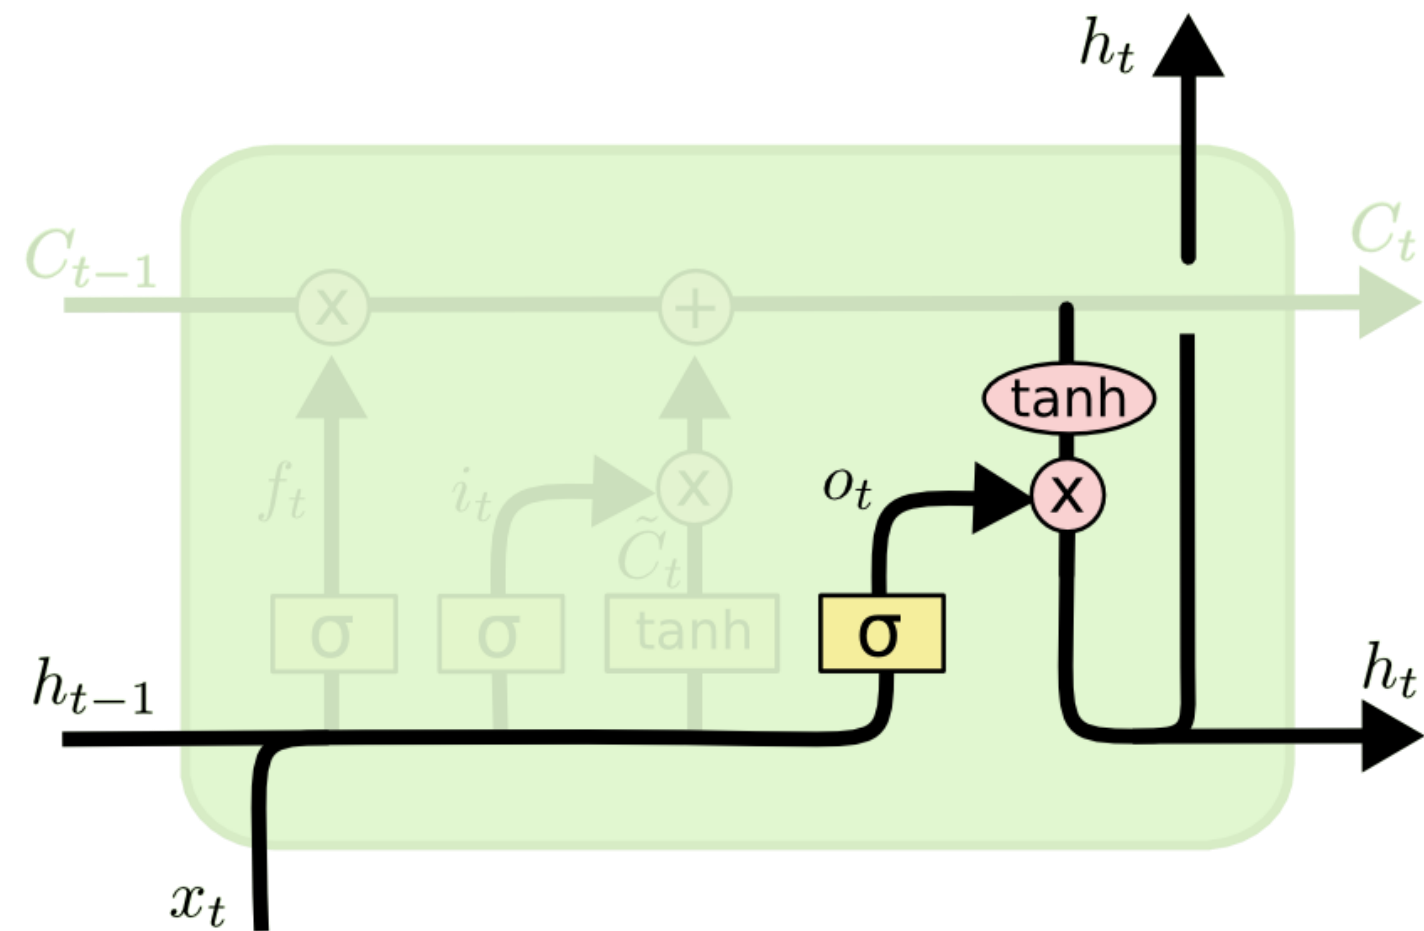
add the gated candidate values

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$$


forget part of the previous cell state



The **output gate**: what part of the cell state should influence the next output?



$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$
$$h_t = o_t \odot \tanh(C_t)$$



Basic structure of an LSTM:

vanilla RNN:

$$h_t = \tanh \left(W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} \right)$$

LSTM:

input gate
forget gate
output gate
candidate update

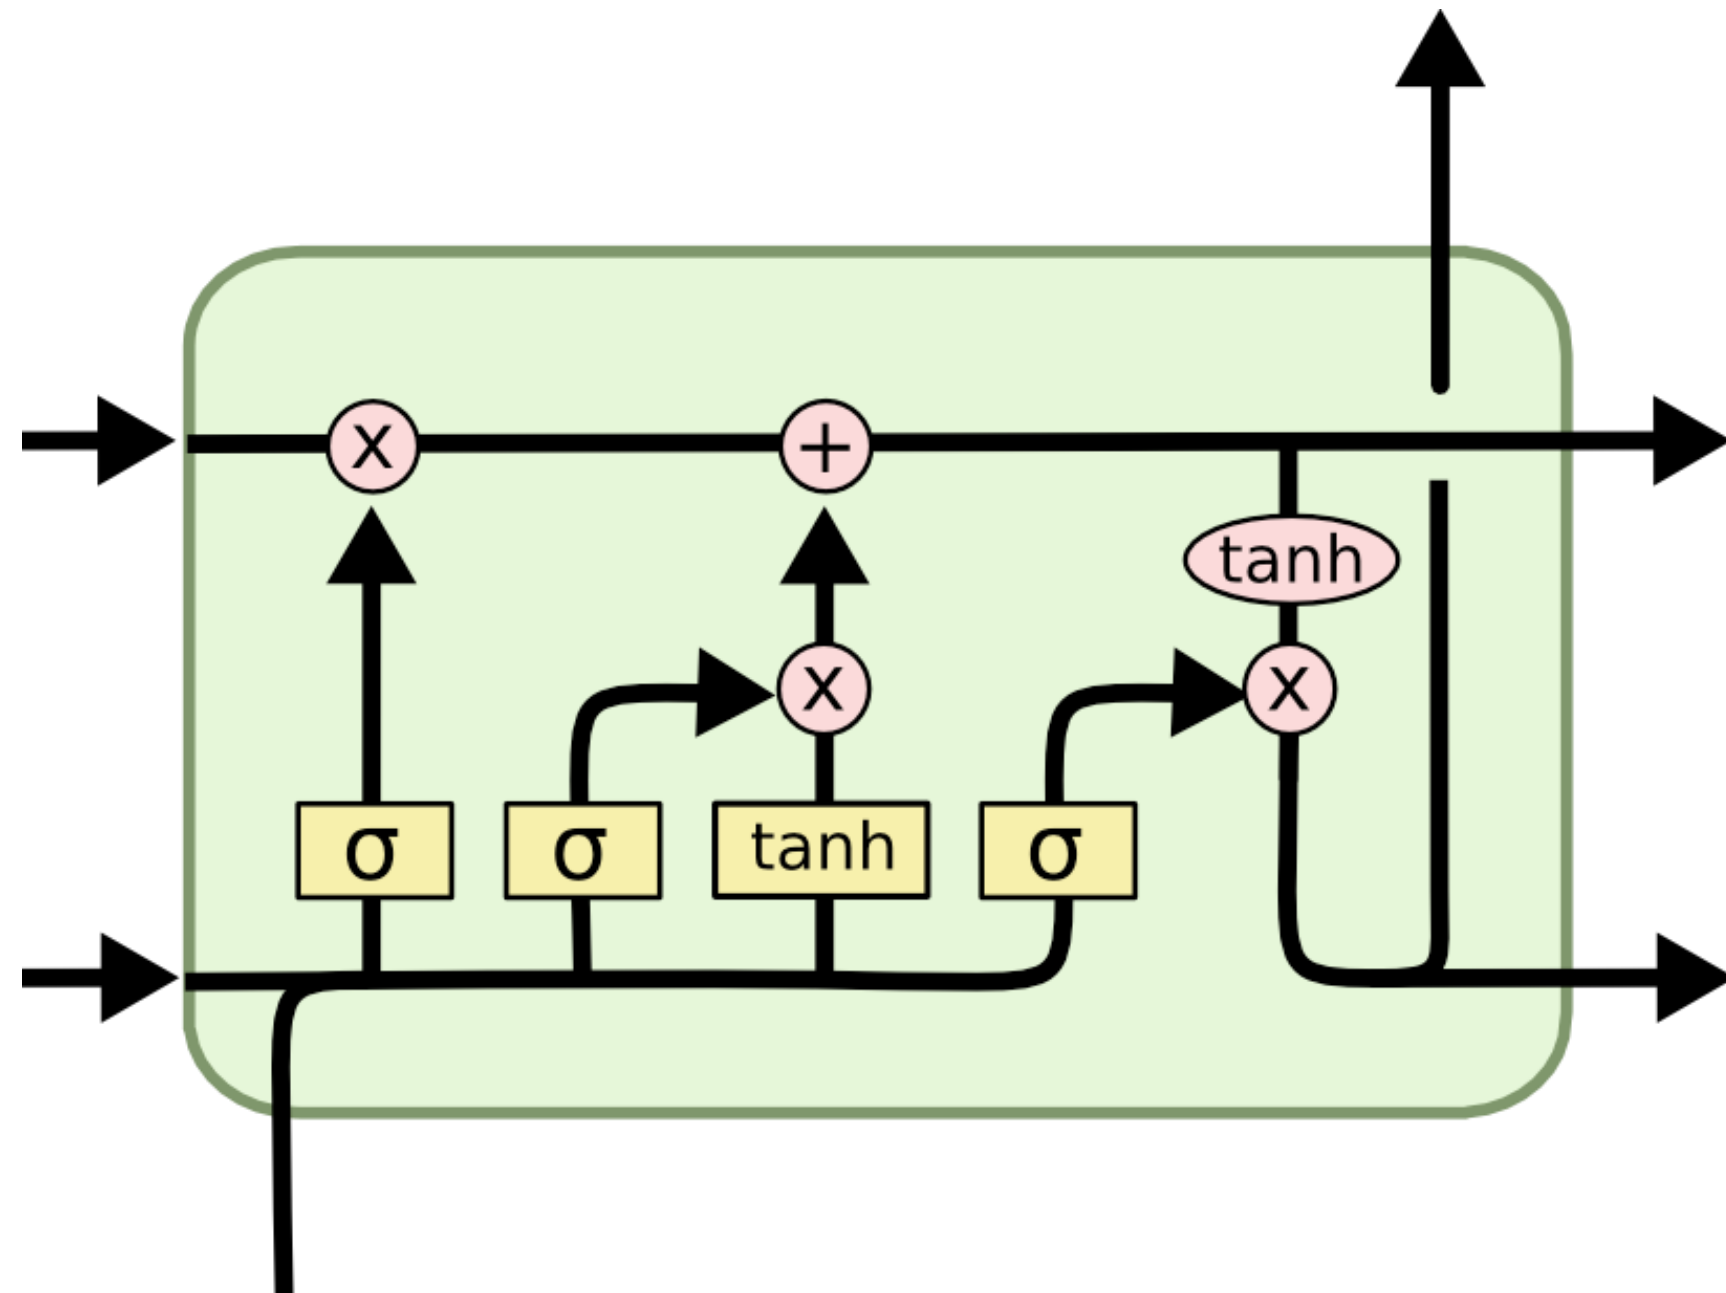
$$\begin{pmatrix} i \\ f \\ o \\ \tilde{c} \end{pmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$

$$c_t = f \odot c_{t-1} + i \odot \tilde{c}$$

$$h_t = o \odot \tanh(c_t)$$

↑
Hadamard product (element-wise multiplication)

Basic structure of an LSTM:

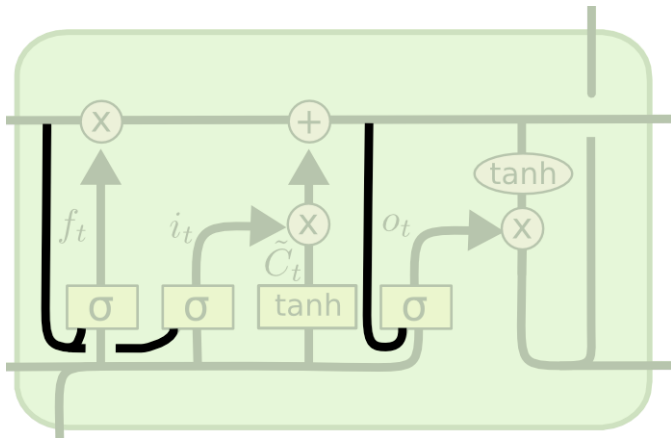


$$\begin{pmatrix} i \\ f \\ o \\ \tilde{c} \end{pmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$

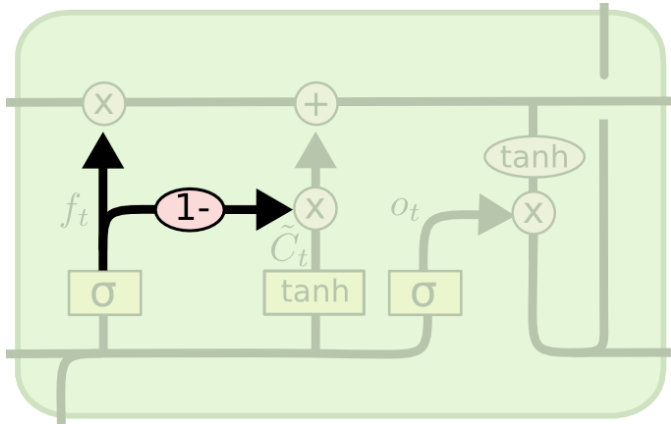
$$c_t = f \odot c_{t-1} + i \odot \tilde{c}$$

$$h_t = o \odot \tanh(c_t)$$

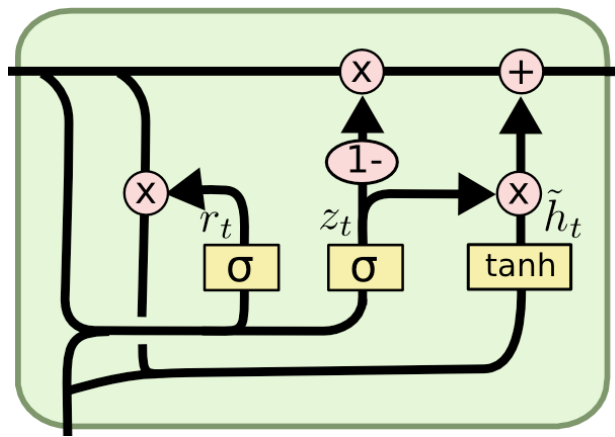
Variations on LSTMs:




peep-hole connections




coupled forget and input gates





GRU (Gated Recurrent Unit),
combines forget and input gates into
a single update gate




Basic structure of a convolutional LSTM:



The core idea behind convolutional LSTMs is the same:




The **forget gate**: is there any information in the cell state that we no longer need?




convolution operation


$$f_t = \sigma(W_f * [h_{t-1}, x_t] + b_f)$$



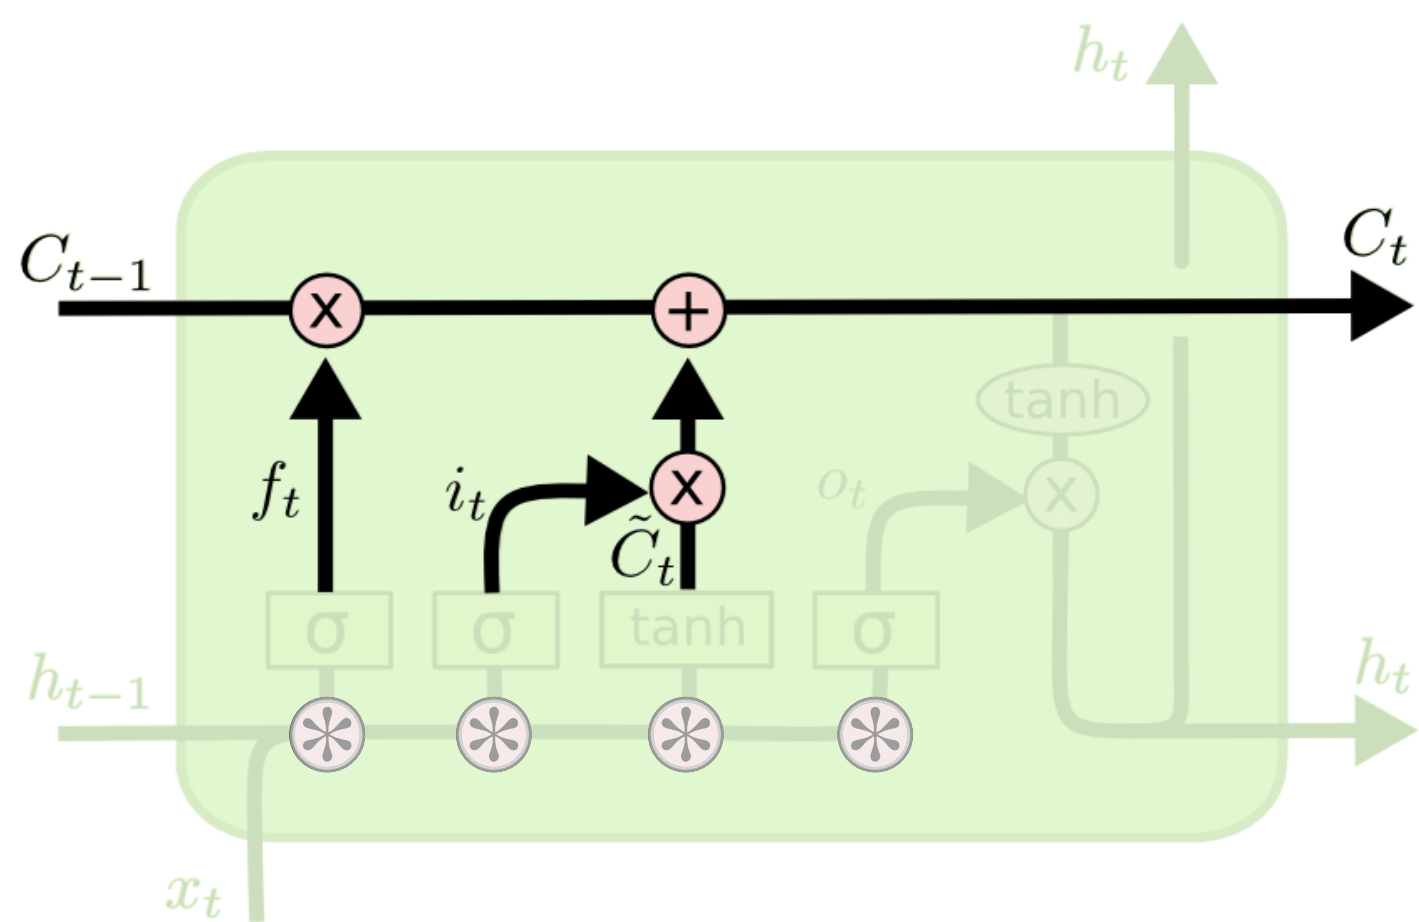
The **input gate** and the **candidate update**: what part of the input should we use to update the cell state?



$$i_t = \sigma(W_i * [h_{t-1}, x_t] + b_i)$$
$$\tilde{C}_t = \tanh(W_c * [h_{t-1}, x_t] + b_c)$$




Updating the cell state:



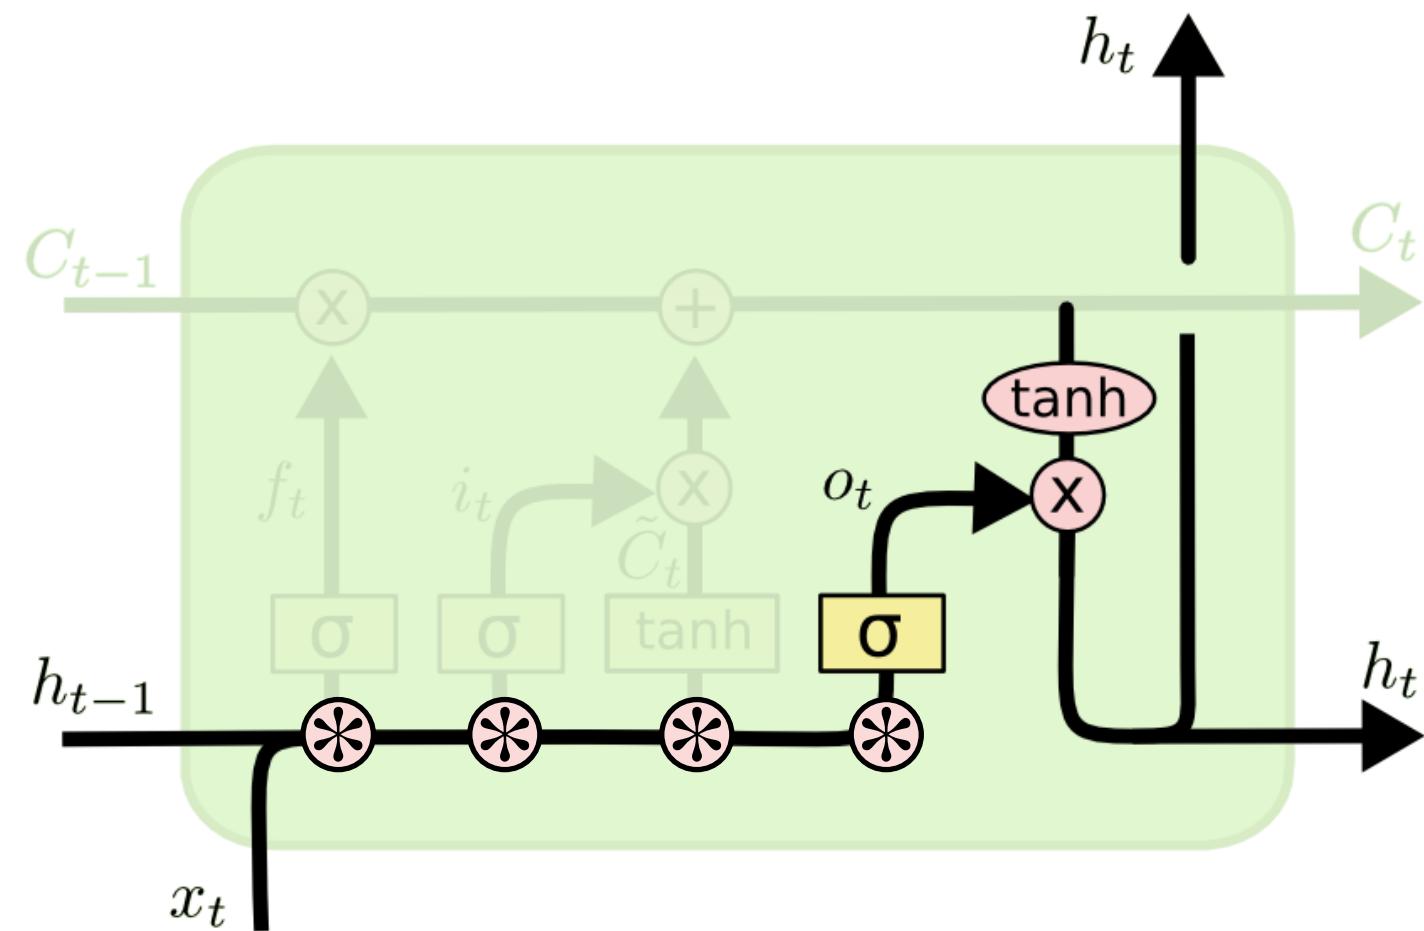
add the gated candidate values

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$$


forget part of the previous cell state



The **output gate**: what part of the cell state should influence the next output?



$$o_t = \sigma(W_o * [h_{t-1}, x_t] + b_o)$$
$$h_t = o_t \odot \tanh(C_t)$$



Basic structure of a convolutional LSTM:

$$f_t = \sigma(W_f * [h_{t-1}, x_t] + b_f)$$

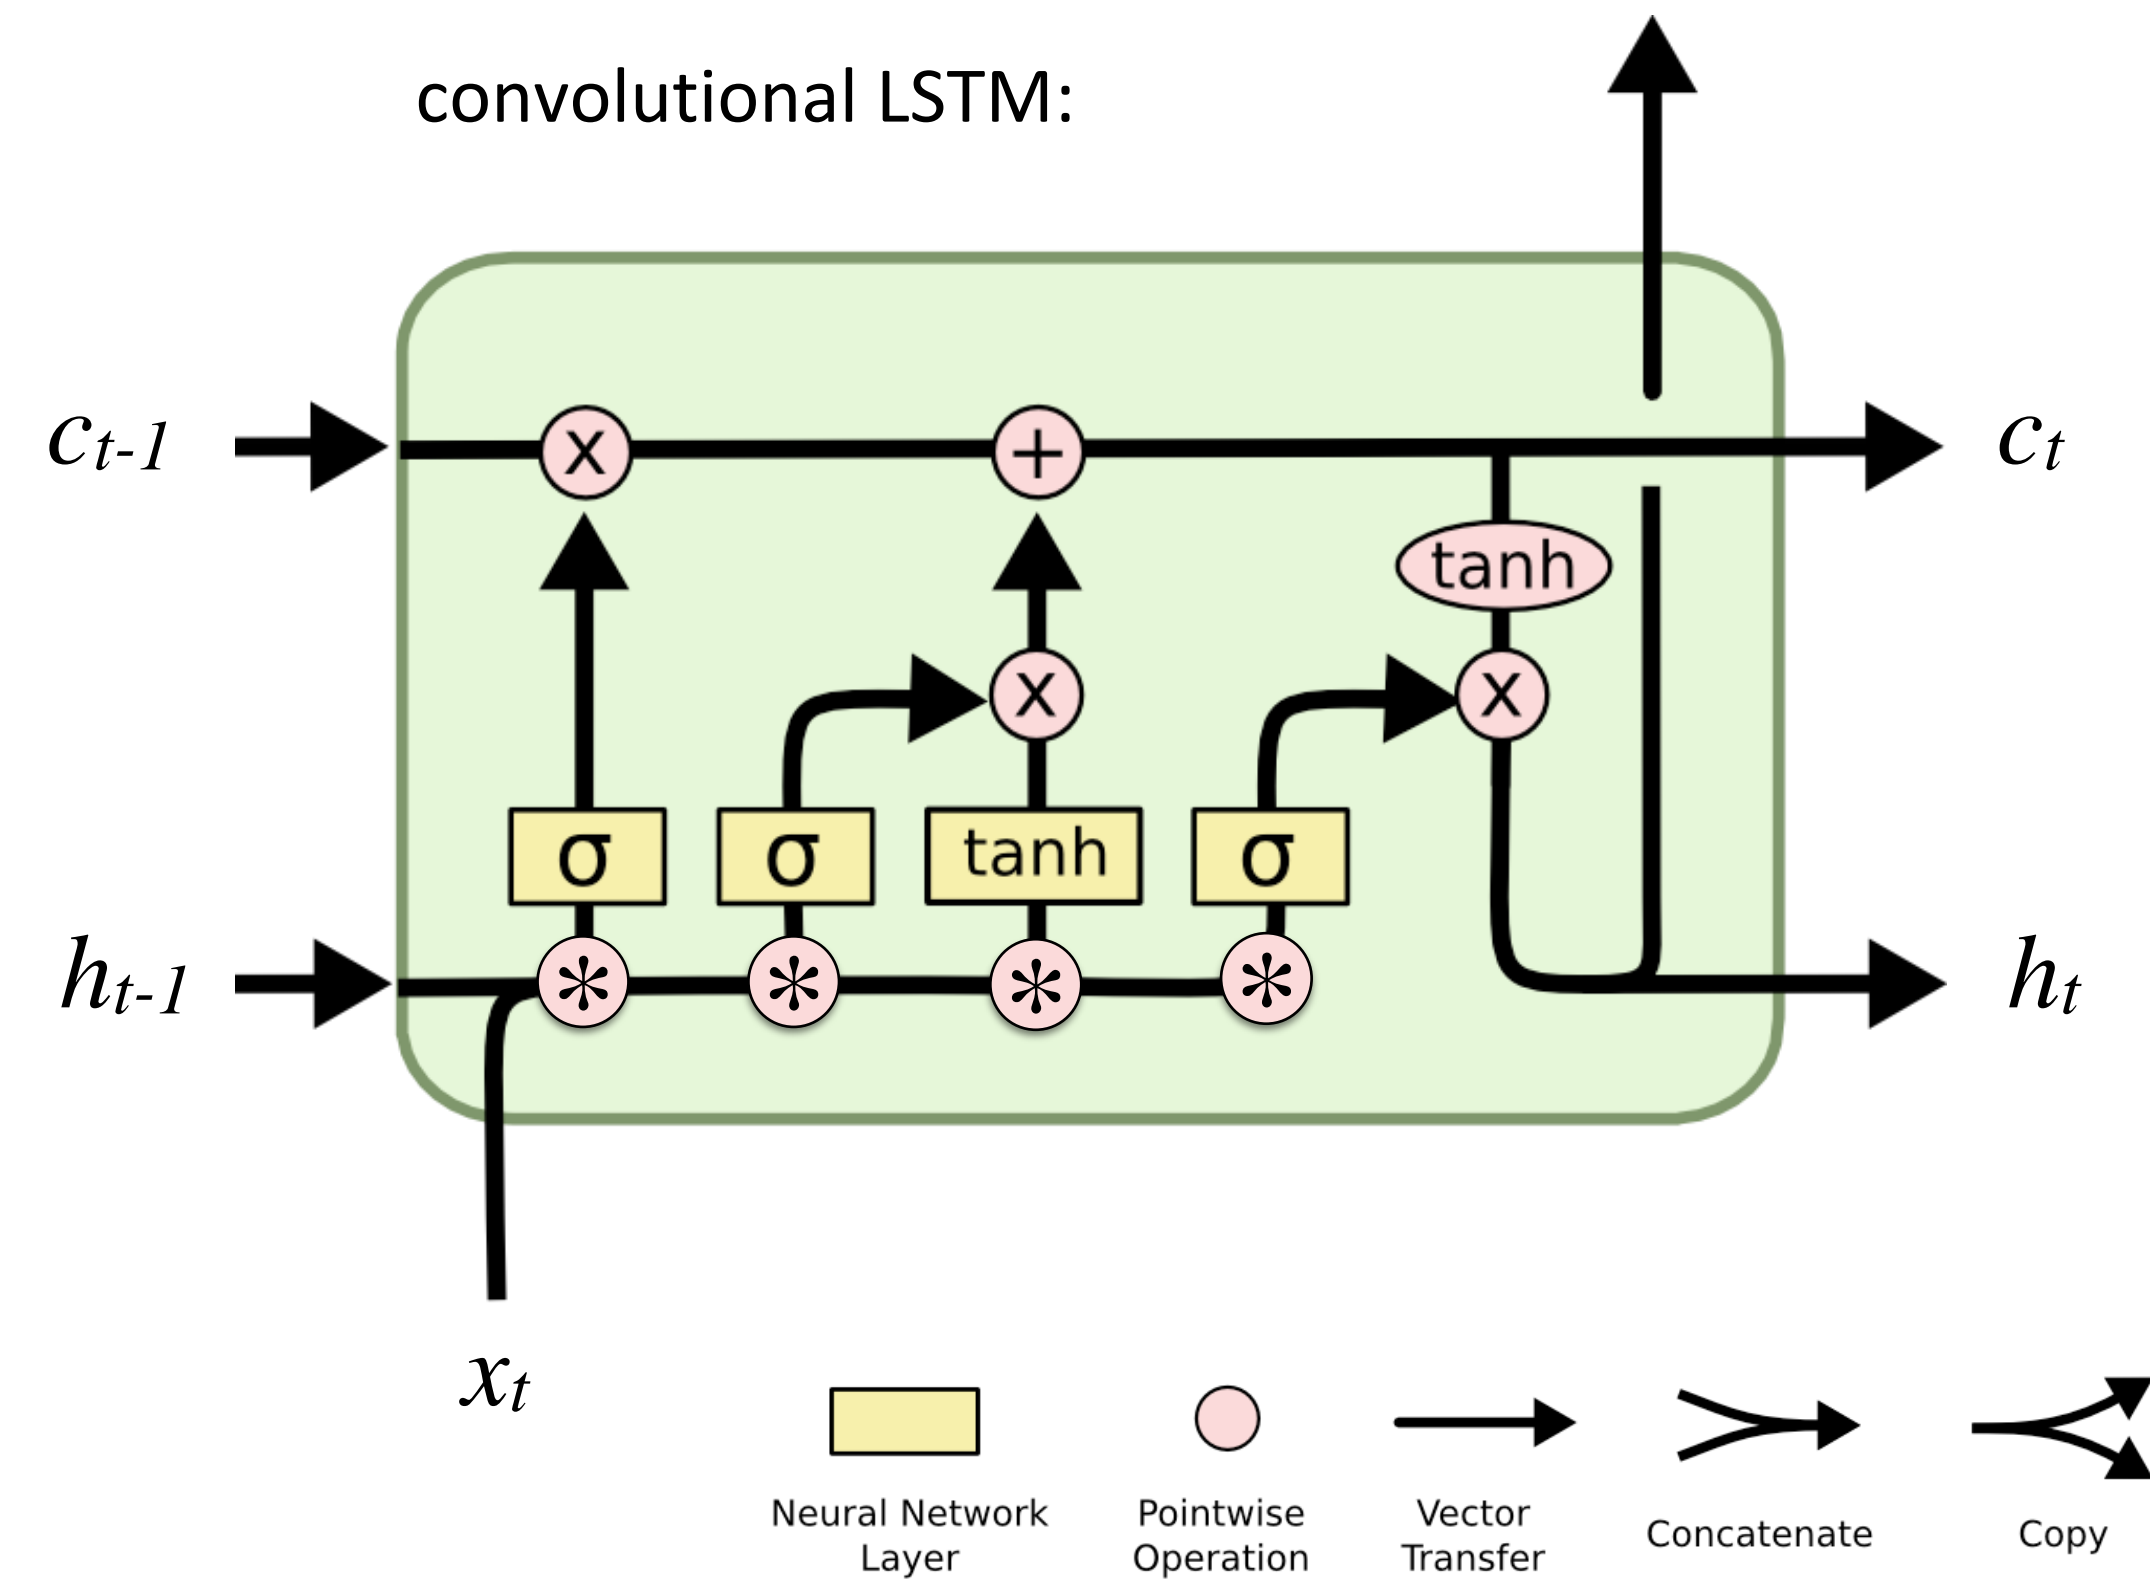
$$i_t = \sigma(W_i * [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_c * [h_{t-1}, x_t] + b_c)$$

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$$

$$o_t = \sigma(W_o * [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t \odot \tanh(C_t)$$



Basic structure of a convolutional LSTM (with peep-hole connections):

$$f_t = \sigma(W_f * [h_{t-1}, x_t] + W_{cf} \odot C_{t-1} + b_f)$$


$$i_t = \sigma(W_i * [h_{t-1}, x_t] + W_{ci} \odot C_{t-1} + b_i)$$

$$\tilde{C}_t = \tanh(W_c * [h_{t-1}, x_t] + b_c)$$


$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$$

$$o_t = \sigma(W_o * [h_{t-1}, x_t] + W_{co} \odot C_{t-1} + b_o)$$

$$h_t = o_t \odot \tanh(C_t)$$



Examples of application:



https://github.com/pytorch/examples/tree/master/time_sequence_prediction

VIOLA:

Why, Salisbury must find his flesh and thought
That which I am not aps, not a man and in fire,
To show the reining of the raven and the wars
To grace my hand reproach within, and not a fair are hand,
That Caesar and my goodly father's world;
When I was heaven of presence and our fleets,
We spare with hours, but cut thy council I am great,
Murdered and by thy master's ready there
My power to give thee but so much as hell:
Some service in the noble bondman here,
Would show him to her wine.

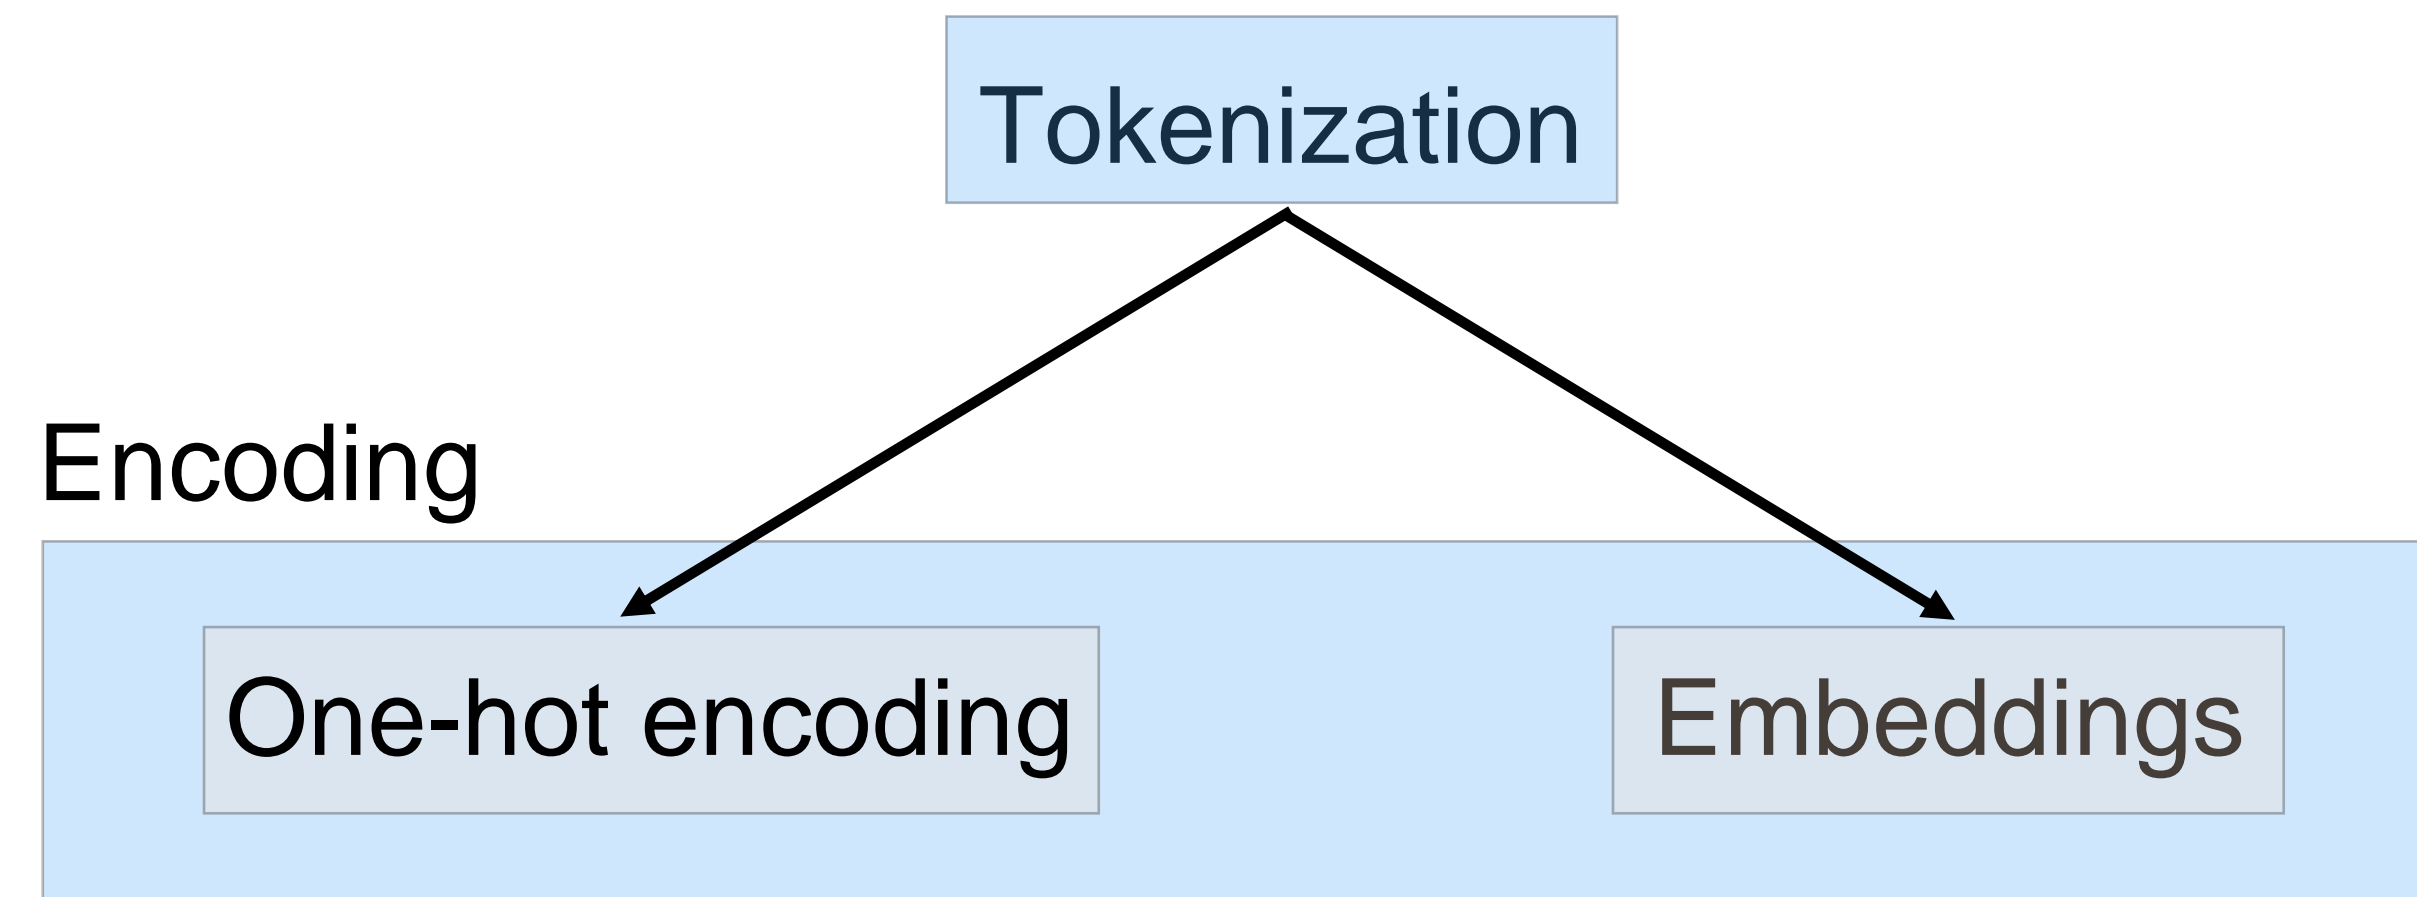
KING LEAR:

O, if you were a feeble sight, the courtesy of your law,
Your sight and several breath, will wear the gods
With his heads, and my hands are wonder'd at the deeds,
So drop upon your lordship's head, and your opinion
Shall be against your honour.

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

1. Recurrent Neural Networks (RNNs)
2. Long Short-Term Memory (LSTM) networks
- 3. Text representation**

Text datasets need to be transformed into numerical form to be used by our networks:

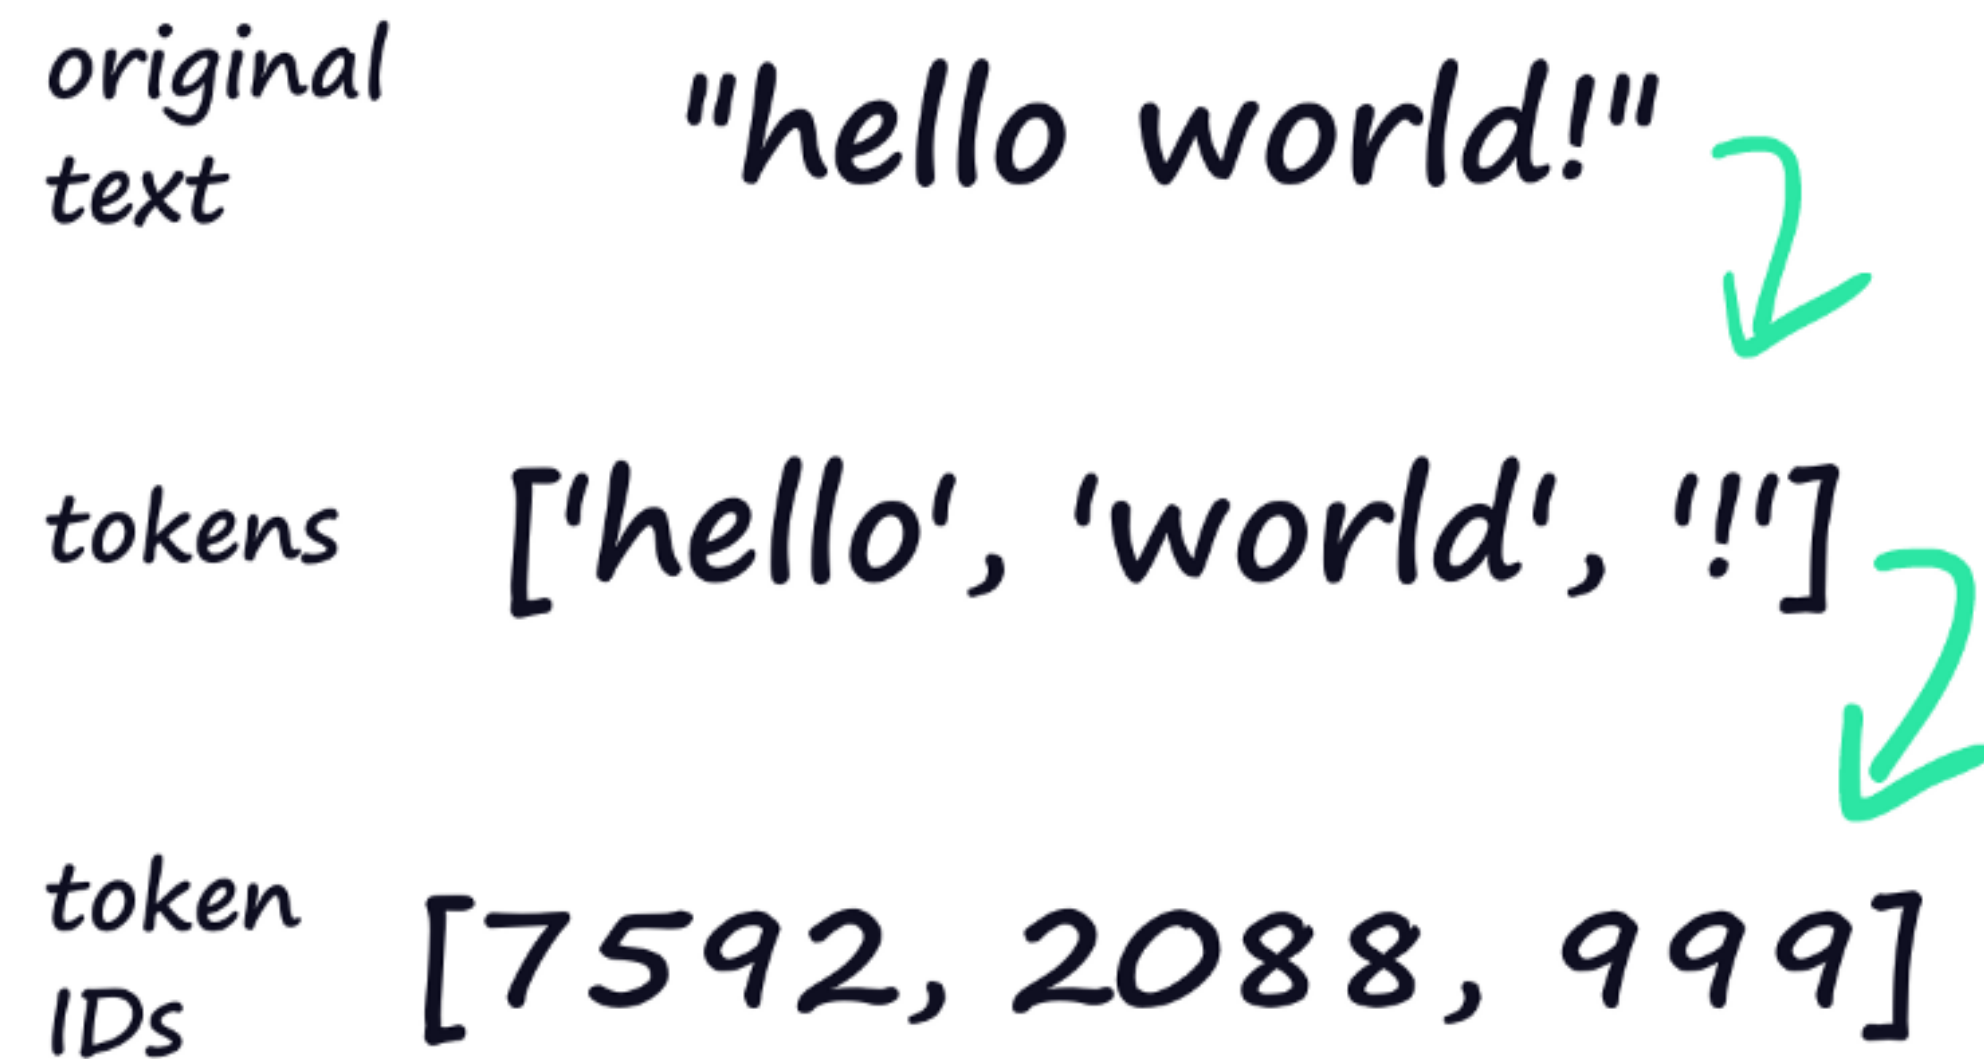


Tokenization is the process of encoding a string of text into numerical integer IDs:

original
text "hello world!"

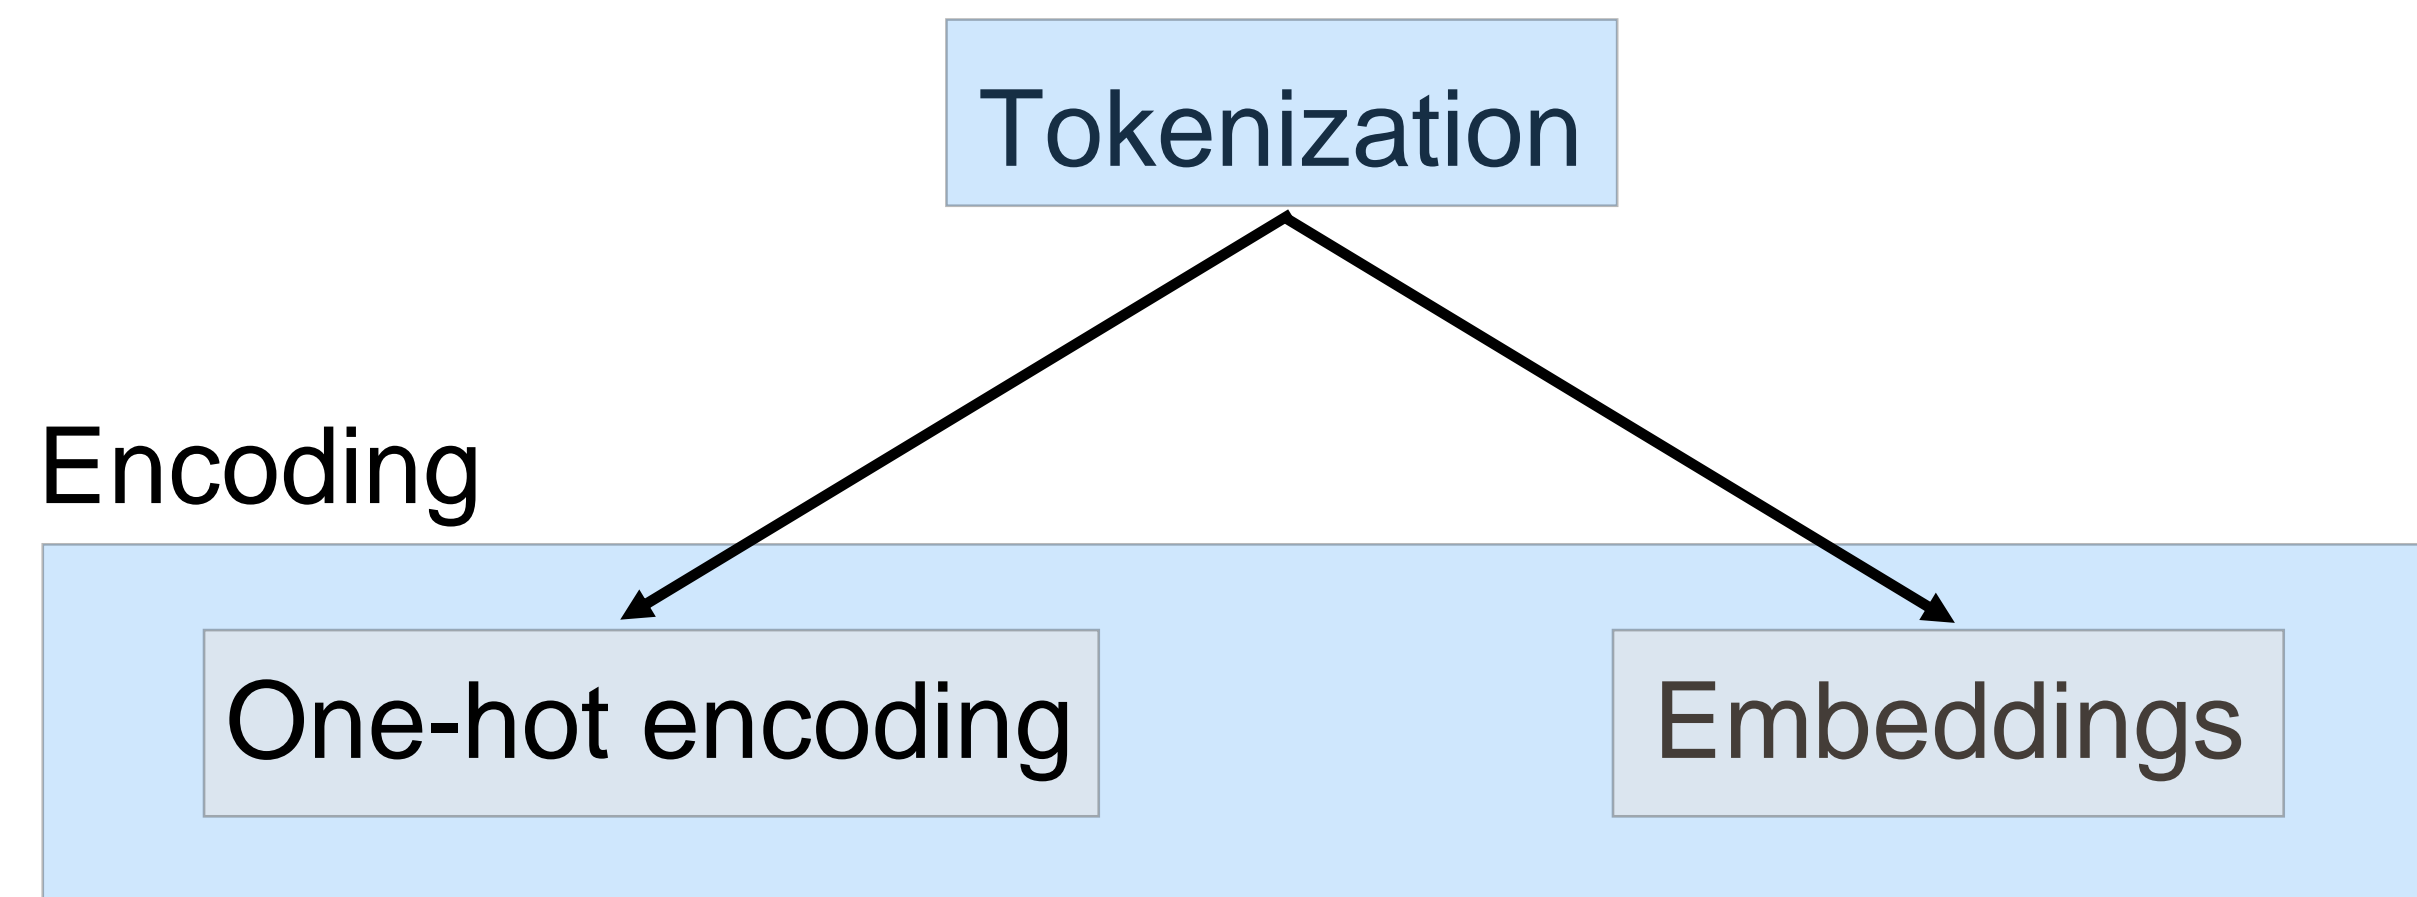
tokens ['hello', 'world', '!']

token
IDs [7592, 2088, 999]

A diagram illustrating the tokenization process. It shows three stages: 'original text' with the string '"hello world!"', 'tokens' with the list ['hello', 'world', '!'], and 'token IDs' with the list [7592, 2088, 999]. Green arrows indicate the flow from the original text to the tokens, and from the tokens to the token IDs.

Text datasets need to be transformed into numerical form to be used by our networks:


Once tokenized, we need to encode tokens into vectors:



With hot encoding we transform numerical integer IDs into vectors of 1s and 0s:


token
IDs *[7592, 2088, 999]*

0	0	0
1	0	0
0	0	0
⋮	⋮	⋮
0	1	0
0	0	0
0	0	1



With hot encoding we transform numerical integer IDs into vectors of 1s and 0s:

token
IDs [7592, 2088, 999]

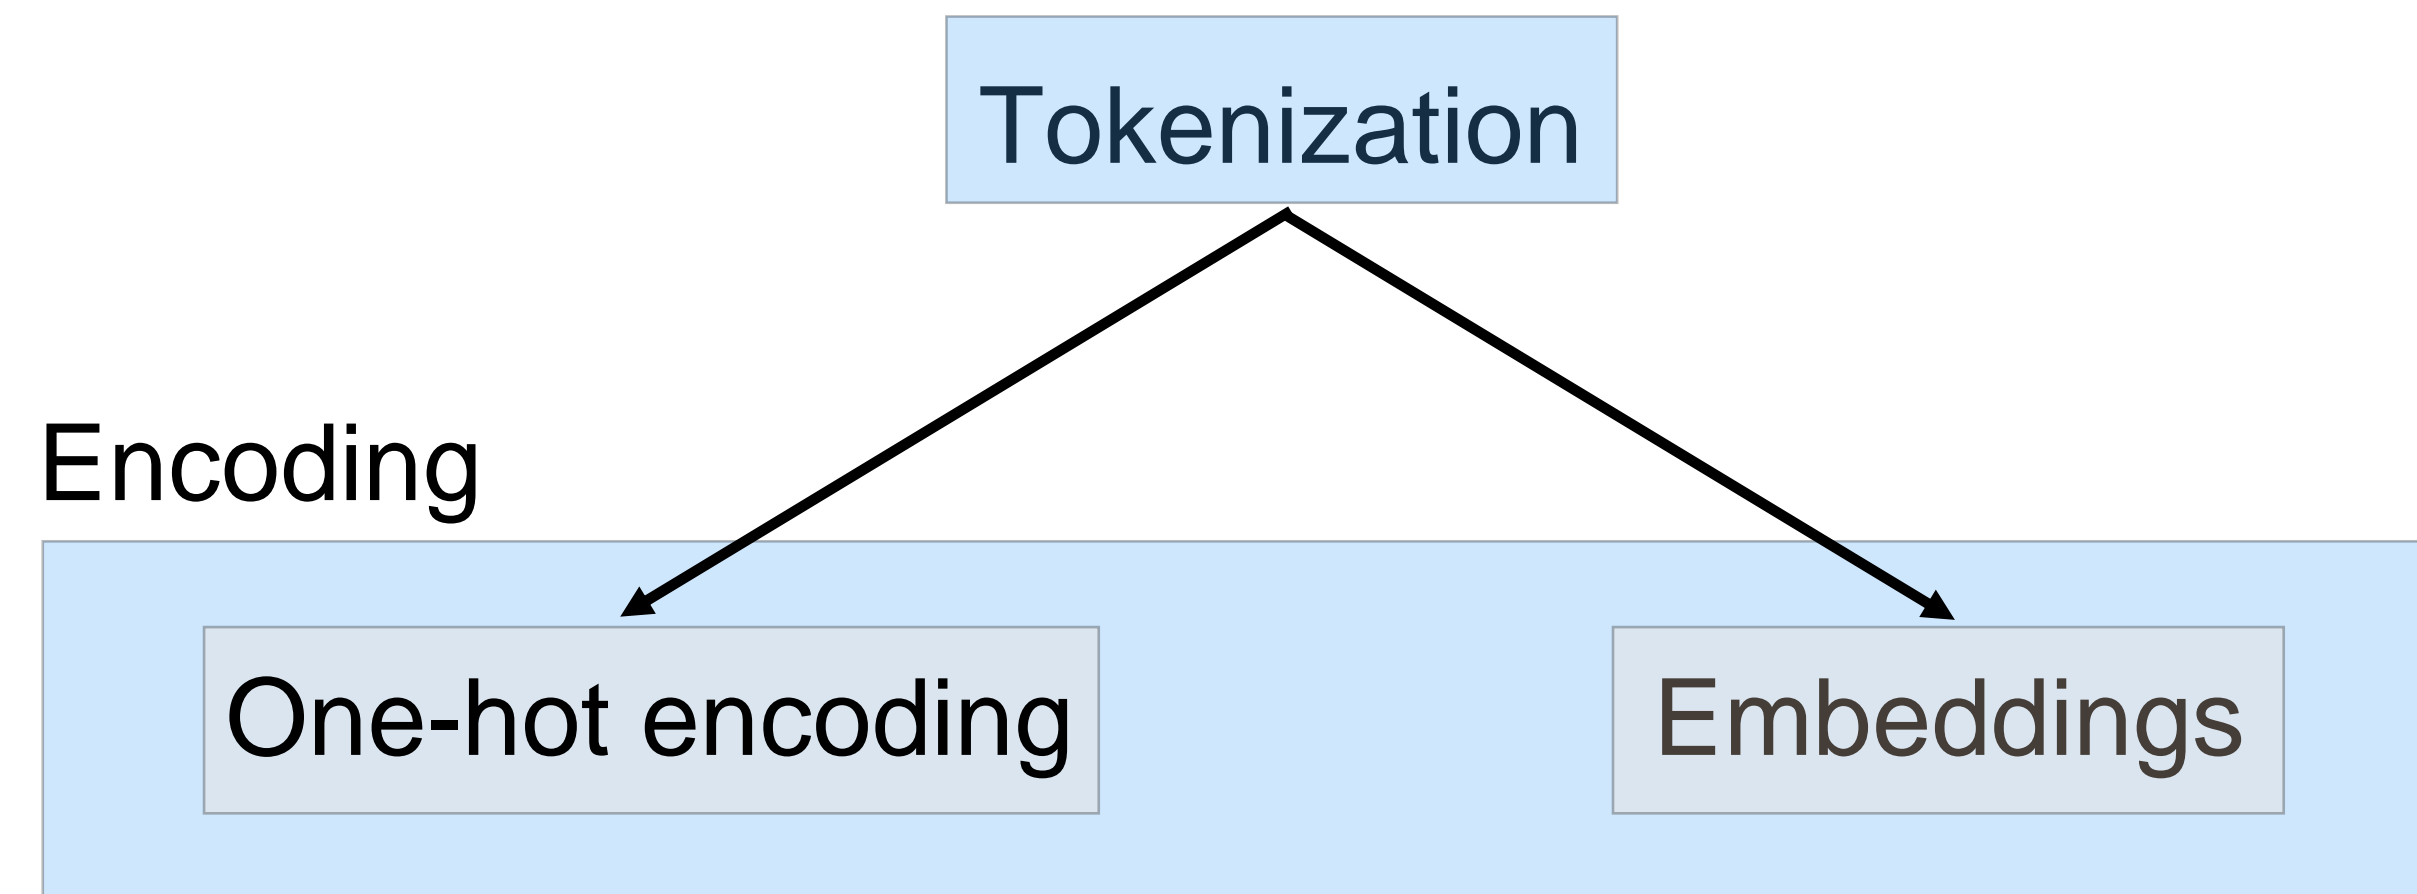


0	0	0
1	0	0
0	0	0
⋮	⋮	⋮
0	1	0
0	0	0
0	0	1


When using words instead of characters, the dictionary would be huge!

Text datasets need to be transformed into numerical form to be used by our networks:


To deal with large dictionaries, we need smaller but representation denser vectors:




Words have meaning and relations, so we can represent word tokens in a dense vector space:



Male-Female

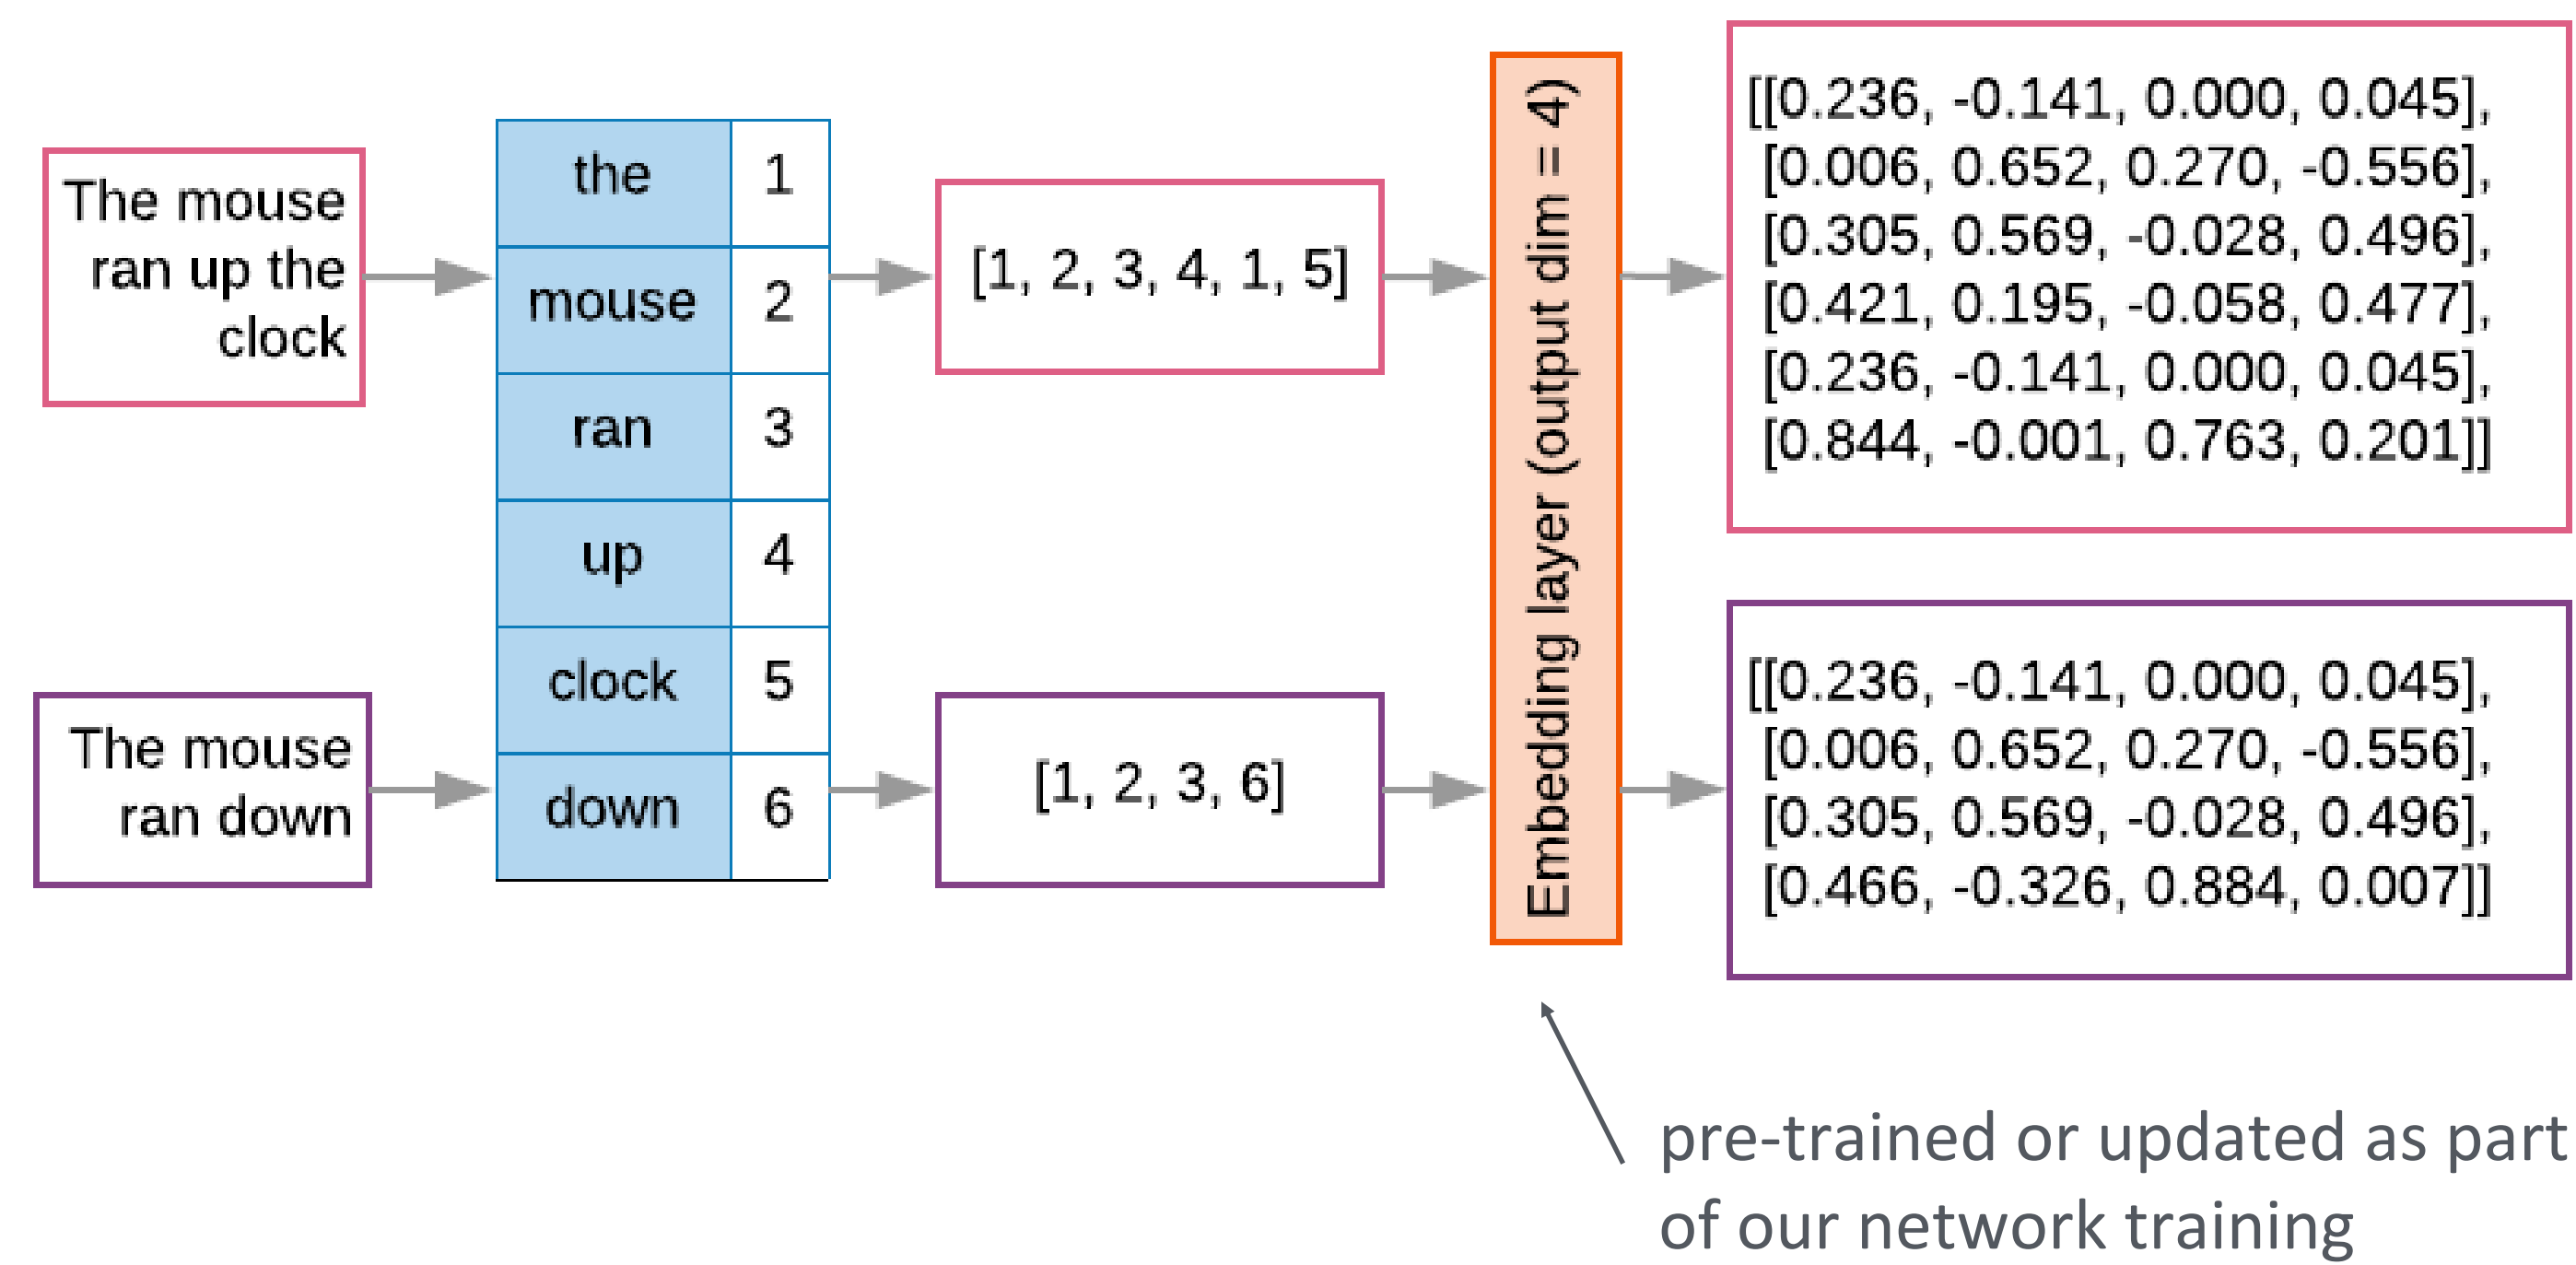


Verb Tense



Country-Capital

Words have meaning and relations, so we can represent word tokens in a dense vector space:



The basis of successful natural language processing (NLP) models like GPT-3 or BERT.

- ▶ Recurrent Neural Networks (RNNs) are the basic approach for sequential data.
- ▶ RNNs suffer from exploding and vanishing gradients and only preserve short-term memory.
- ▶ Long Short-Term Memory (LSTM) networks address these two issues.
- ▶ Before text can be processed by neural networks it needs to be tokenized into unique IDs.
- ▶ Tokenized text can be encoded using one-hot encoding (appropriate for small dictionaries) or using embeddings (more efficient for larger dictionaries).